

TIME SERIES FORECASTING

PROJECT

Abstract

**the data of different types of wine sales in the 20th century is to be analyzed.
Both of these data are from the same company but of different wines. As an
analyst in the ABC Estate Wines, you are tasked to analyses and forecast Wine
Sales in the 20th century.**

Rekha Wankhede
[Email address]

1. Read the data as an appropriate Time Series data and plot the data.

- Both Datasets are read and stored as Pandas Data Frames for analysis
 - Datasets are read as Time Series data using `parse dates=True` & `index_col='YearMonth'`
 - First 5 rows of both the data are given below .

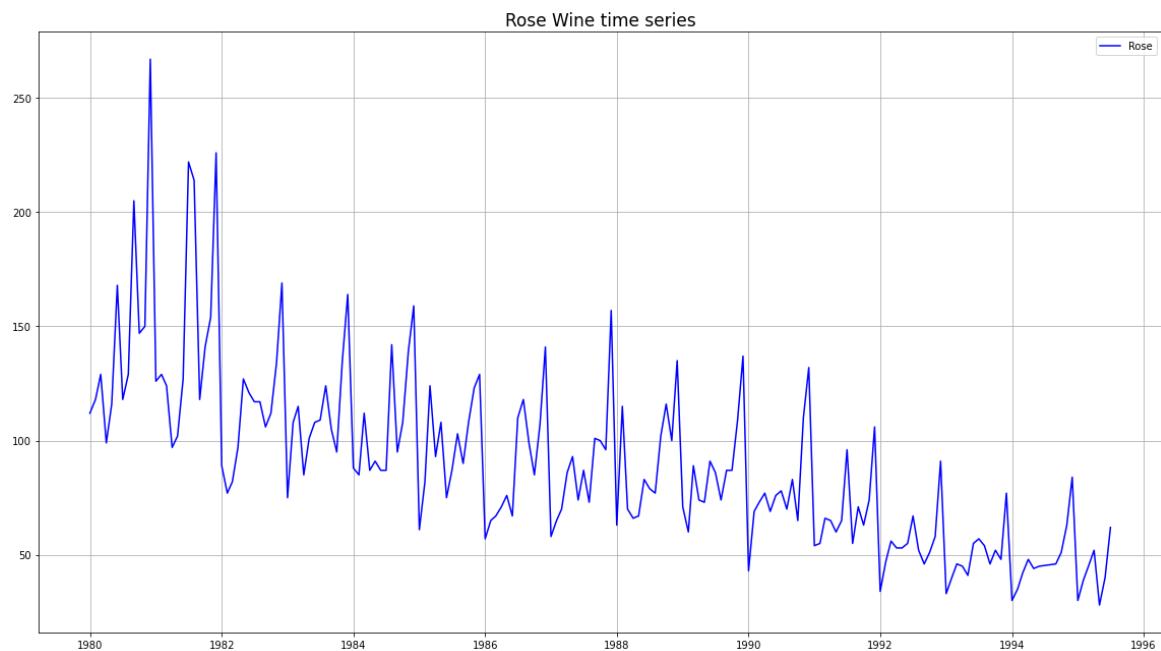
Rose Data Header

YearMonth	Rose
1980-01-01	112.0
1980-02-01	118.0
1980-03-01	129.0
1980-04-01	99.0
1980-05-01	116.0

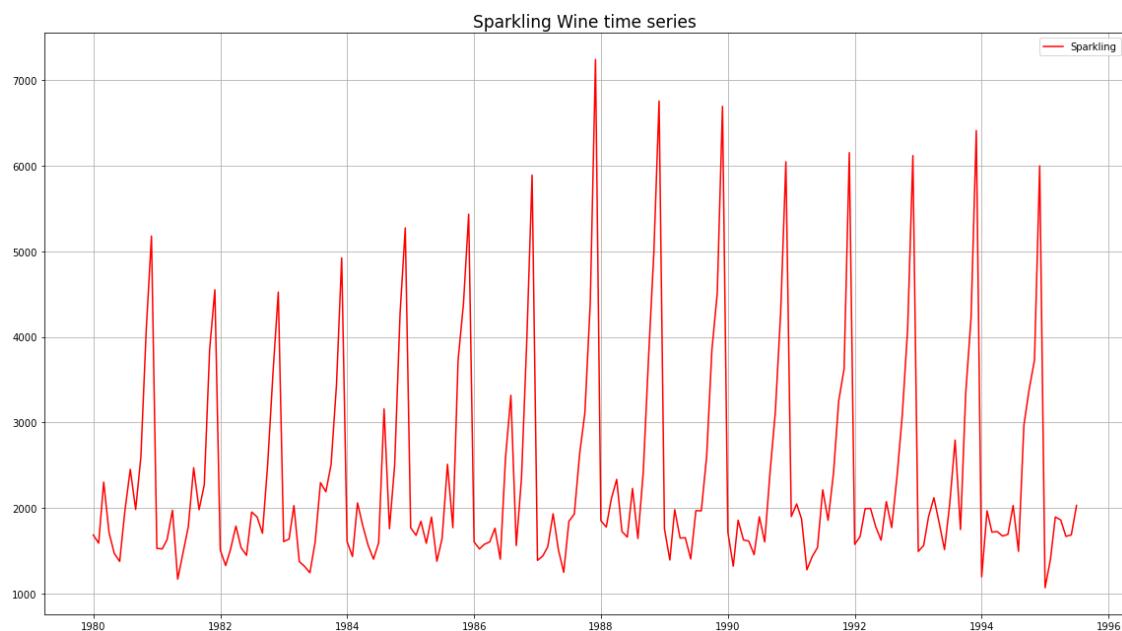
Sparkling Data Header

YearMonth	Sparkling
1980-01-01	1686
1980-02-01	1591
1980-03-01	2304
1980-04-01	1712
1980-05-01	1471

- Rose Data plot



• **Sparkling Data plot**



2. Perform appropriate Exploratory Data Analysis to understand the data and also perform decomposition.

Exploratory Data Analysis:

Descriptive Stats of Rose and Sparkling datasets.

	count	mean	std	min	25%	50%	75%	max
Rose	187.0	89.914439	39.238325	28.0	62.5	85.0	111.0	267.0

	count	mean	std	min	25%	50%	75%	max
Sparkling	187.0	2402.417112	1295.11154	1070.0	1605.0	1874.0	2549.0	7242.0

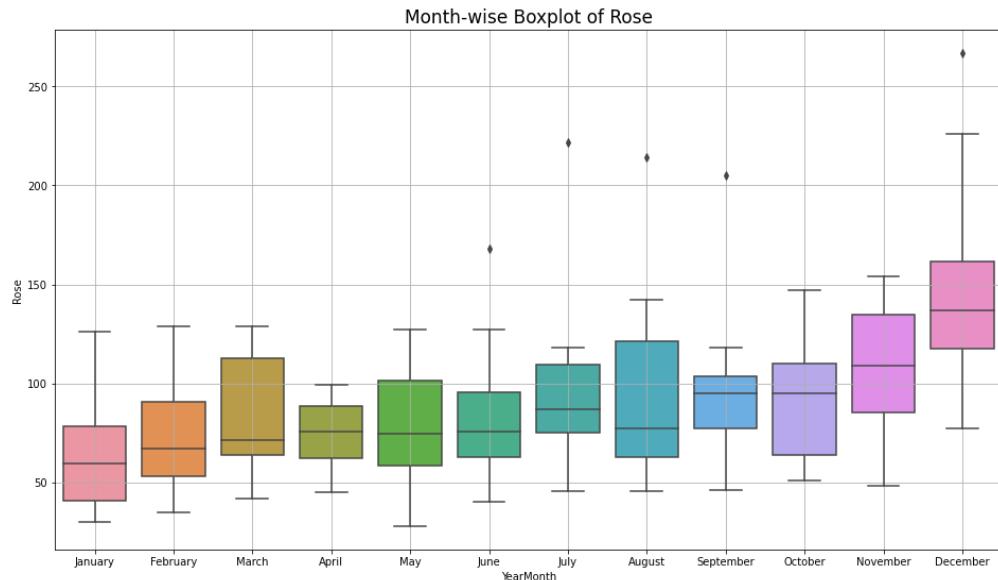
```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 187 entries, 1980-01-01 to 1995-07-01
Data columns (total 1 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Rose        185 non-null    float64
dtypes: float64(1)
memory usage: 2.9 KB
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 187 entries, 1980-01-01 to 1995-07-01
Data columns (total 1 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Sparkling   187 non-null    int64 
dtypes: int64(1)
memory usage: 2.9 KB
```

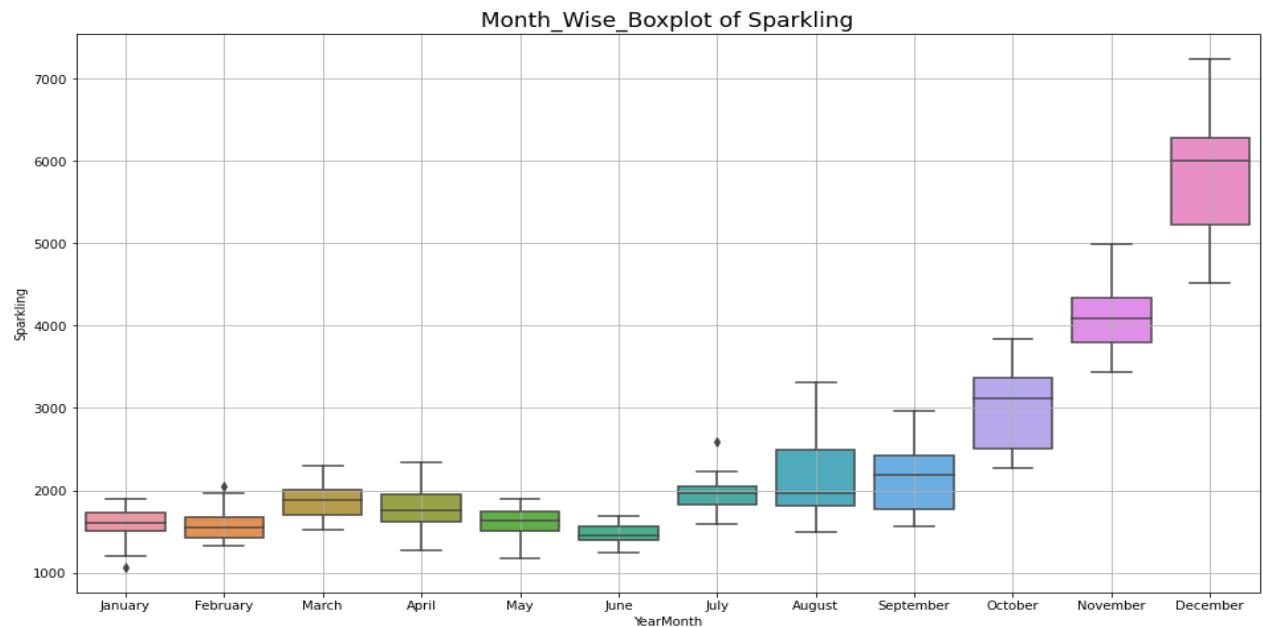
ROSE data. info

Sparkling DATASET INFO

- Month-wise Boxplot of Rose

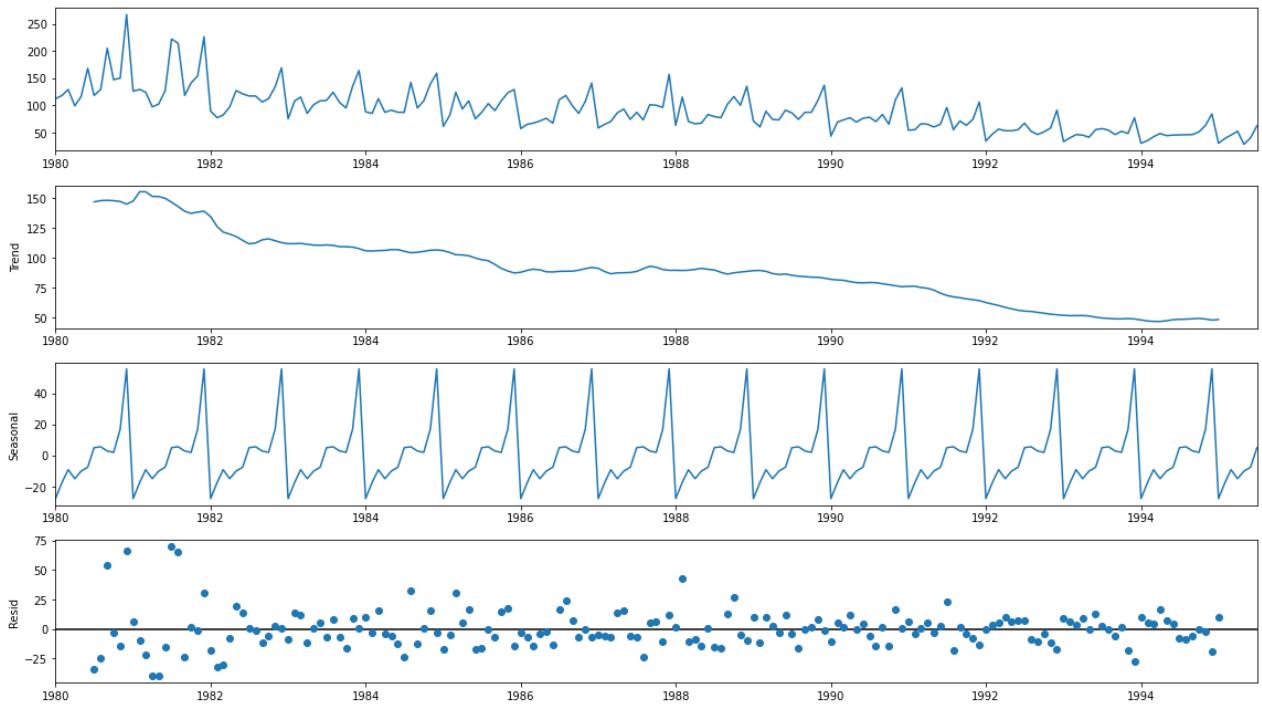


- Month-wise Boxplot of Sparkling -



- Sales of both - Rose and Sparkling, show a spike in the last quarter of Oct to Dec
 - Spike is much more accentuated in Sparkling sales
 - This spike may be due to the Holiday season starting in Oct

Additive Decomposition of Rose –



```

Trend
YearMonth
1980-01-01      NaN
1980-02-01      NaN
1980-03-01      NaN
1980-04-01      NaN
1980-05-01      NaN
1980-06-01      NaN
1980-07-01  147.083333
1980-08-01  148.125000
1980-09-01  148.375000
1980-10-01  148.083333
1980-11-01  147.416667
1980-12-01  145.125000
Name: trend, dtype: float64
  
```

```

Seasonality
YearMonth
1980-01-01  -27.908647
1980-02-01  -17.435632
1980-03-01  -9.285830
1980-04-01  -15.098330
1980-05-01  -10.196544
1980-06-01  -7.678687
1980-07-01   4.896908
1980-08-01   5.499686
1980-09-01   2.774686
1980-10-01   1.871908
1980-11-01  16.846908
1980-12-01  55.713575
Name: seasonal, dtype: float64
  
```

```

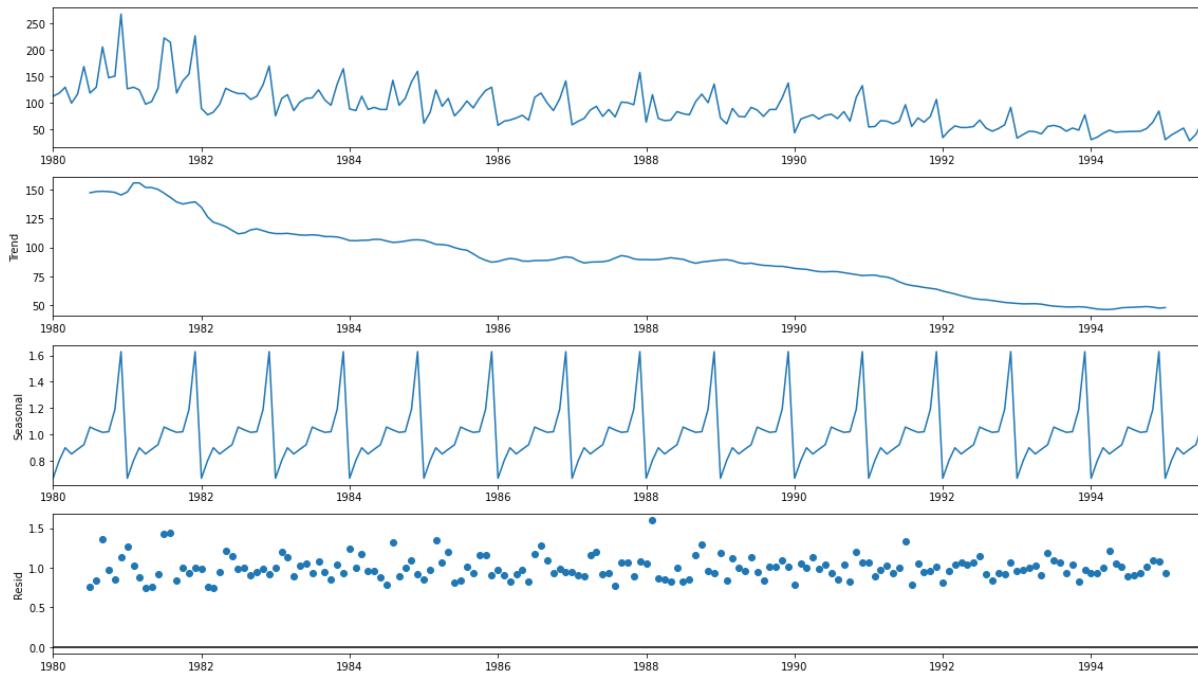
Residual
YearMonth
1980-01-01      NaN
1980-02-01      NaN
1980-03-01      NaN
1980-04-01      NaN
1980-05-01      NaN
1980-06-01      NaN
1980-07-01  -33.980241
1980-08-01  -24.624686
1980-09-01  53.850314
1980-10-01  -2.955241
1980-11-01  -14.263575
1980-12-01  66.161425
Name: resid, dtype: float64
  
```

ROSE-TREND

SEASONALITY

Residual

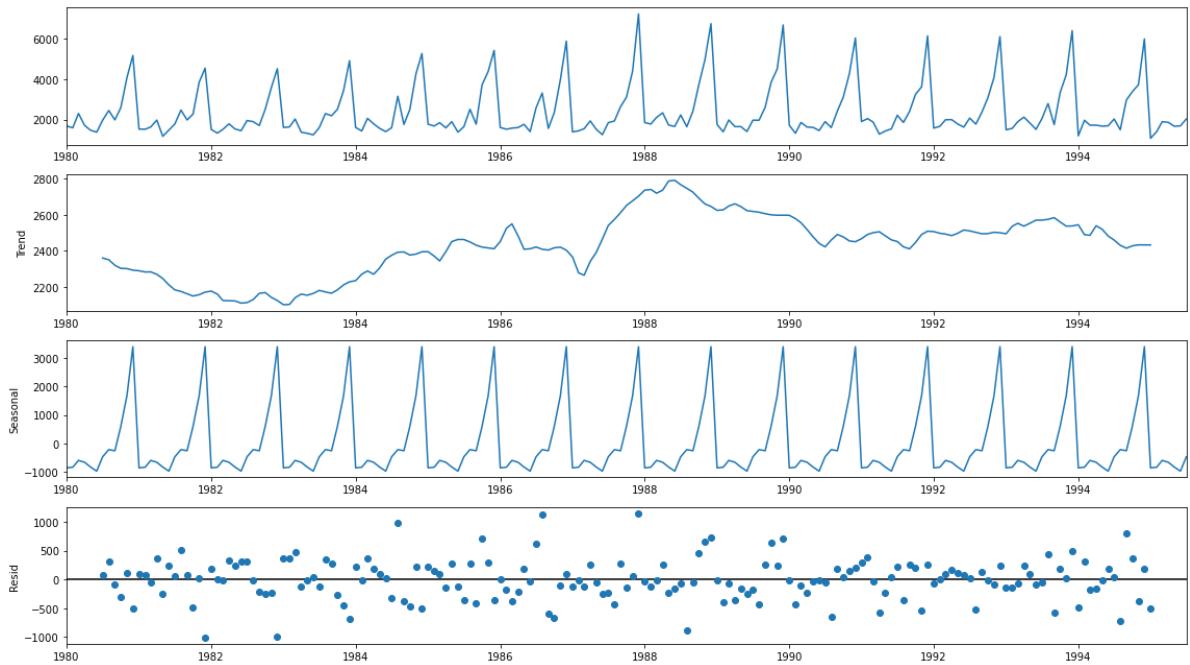
Multiplicative Decomposition of Rose –



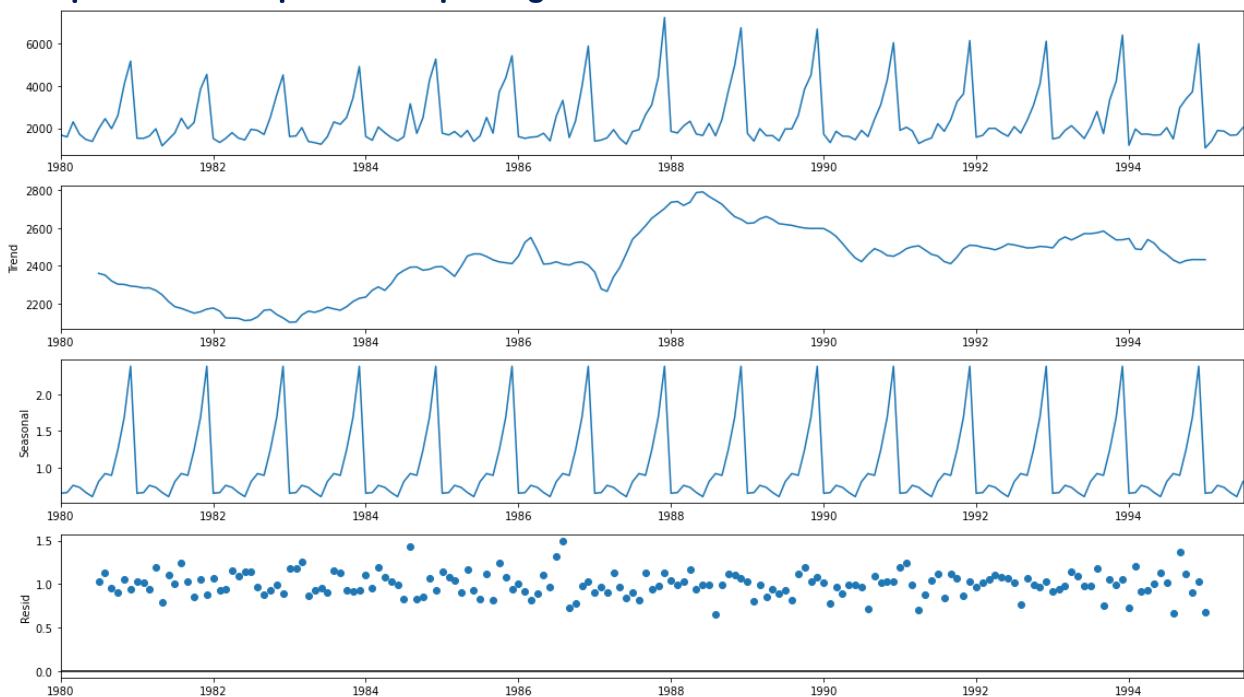
Trend		Seasonality		Residual	
YearMonth		YearMonth		YearMonth	
1980-01-01	NaN	1980-01-01	0.670111	1980-01-01	NaN
1980-02-01	NaN	1980-02-01	0.806163	1980-02-01	NaN
1980-03-01	NaN	1980-03-01	0.901164	1980-03-01	NaN
1980-04-01	NaN	1980-04-01	0.854024	1980-04-01	NaN
1980-05-01	NaN	1980-05-01	0.889415	1980-05-01	NaN
1980-06-01	NaN	1980-06-01	0.923985	1980-06-01	NaN
1980-07-01	147.083333	1980-07-01	1.058038	1980-07-01	0.758258
1980-08-01	148.125000	1980-08-01	1.035881	1980-08-01	0.840720
1980-09-01	148.375000	1980-09-01	1.017648	1980-09-01	1.357674
1980-10-01	148.083333	1980-10-01	1.022573	1980-10-01	0.970771
1980-11-01	147.416667	1980-11-01	1.192349	1980-11-01	0.853378
1980-12-01	145.125000	1980-12-01	1.628646	1980-12-01	1.129646
Name: trend, dtype: float64		Name: seasonal, dtype: float64		Name: resid, dtype: float64	

Additive Decomposition of Sparkling -

Trend		Seasonality		Residual	
YearMonth		YearMonth		YearMonth	
1980-01-01	NaN	1980-01-01	-854.260599	1980-01-01	NaN
1980-02-01	NaN	1980-02-01	-830.350678	1980-02-01	NaN
1980-03-01	NaN	1980-03-01	-592.356630	1980-03-01	NaN
1980-04-01	NaN	1980-04-01	-658.490559	1980-04-01	NaN
1980-05-01	NaN	1980-05-01	-824.416154	1980-05-01	NaN
1980-06-01	NaN	1980-06-01	-967.434011	1980-06-01	NaN
1980-07-01	2360.666667	1980-07-01	-465.502265	1980-07-01	70.835599
1980-08-01	2351.333333	1980-08-01	-214.332821	1980-08-01	315.999487
1980-09-01	2320.541667	1980-09-01	-254.677265	1980-09-01	-81.864401
1980-10-01	2303.583333	1980-10-01	599.769957	1980-10-01	-307.353290
1980-11-01	2302.041667	1980-11-01	1675.067179	1980-11-01	109.891154
1980-12-01	2293.791667	1980-12-01	3386.983846	1980-12-01	-501.775513
Name: trend, dtype: float64		Name: seasonal, dtype: float64		Name: resid, dtype: float64	



Multiplicative Decomposition of Sparkling:



```

Trend
YearMonth
1980-01-01      NaN
1980-02-01      NaN
1980-03-01      NaN
1980-04-01      NaN
1980-05-01      NaN
1980-06-01      NaN
1980-07-01  2360.666667
1980-08-01  2351.333333
1980-09-01  2320.541667
1980-10-01  2303.583333
1980-11-01  2302.041667
1980-12-01  2293.791667
Name: trend, dtype: float64

```

SPARKLING TREND

```

Seasonality
YearMonth
1980-01-01    0.649843
1980-02-01    0.659214
1980-03-01    0.757440
1980-04-01    0.730351
1980-05-01    0.660609
1980-06-01    0.603468
1980-07-01    0.809164
1980-08-01    0.918822
1980-09-01    0.894367
1980-10-01    1.241789
1980-11-01    1.690158
1980-12-01    2.384776
Name: seasonal, dtype: float64

```

SEASONALITY

```

Residual
YearMonth
1980-01-01      NaN
1980-02-01      NaN
1980-03-01      NaN
1980-04-01      NaN
1980-05-01      NaN
1980-06-01      NaN
1980-07-01  1.029230
1980-08-01  1.135407
1980-09-01  0.955954
1980-10-01  0.907513
1980-11-01  1.050423
1980-12-01  0.946770
Name: resid, dtype: float64

```

RESIDUAL

Additive Models –

- Additive Models –

- The seasonality is relatively constant over time

$$Ty = \text{Trend} + \text{Seasonality} y + \text{Residual}$$

Multiplicative Models:

The seasonality increases or decreases over time. It is proportionate to the trend

$$Ty = \text{Trend} * \text{Seasonality} y * \text{Residual}$$

Here by just observing the Residual patterns of Additive and Multiplicative models of Rose and Sparkling datasets. It seems that –

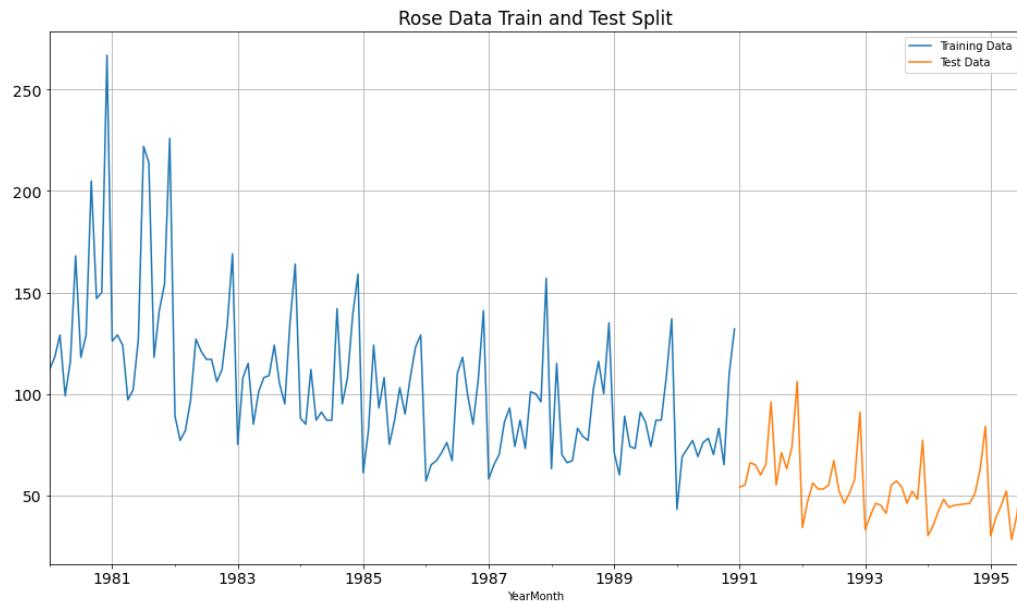
Rose is Multiplicative

Sparkling is Additive

3. **Split the data into training and test. The test data should start in 1991.**

Both datasets of Rose and Sparkling are split at the year 1991

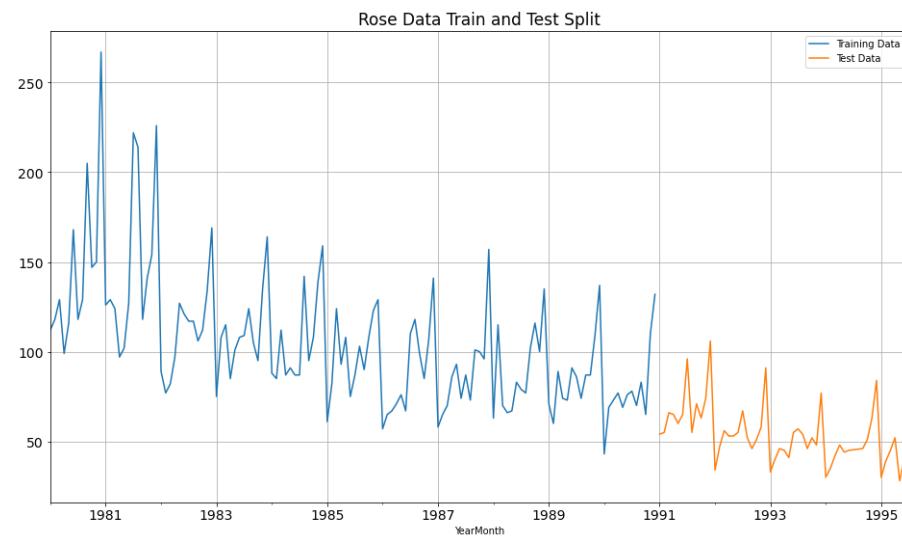
- Test datasets start at 1991



Rose

YearMonth	Rose
1980-01-01	112.0
1980-02-01	118.0
1980-03-01	129.0
1980-04-01	99.0
1980-05-01	116.0
Last few rows of Rose Training Data	
1990-08-01	70.0
1990-09-01	83.0
1990-10-01	65.0
1990-11-01	110.0
1990-12-01	132.0

YearMonth	Rose
1991-01-01	54.0
1991-02-01	55.0
1991-03-01	66.0
1991-04-01	65.0
1991-05-01	60.0
Last few rows of Rose Test Data	
1995-03-01	45.0
1995-04-01	52.0
1995-05-01	28.0
1995-06-01	40.0
1995-07-01	62.0



SPARKLING DATA TRAIN AND TEST

Last few rows of Sparkling Training Data

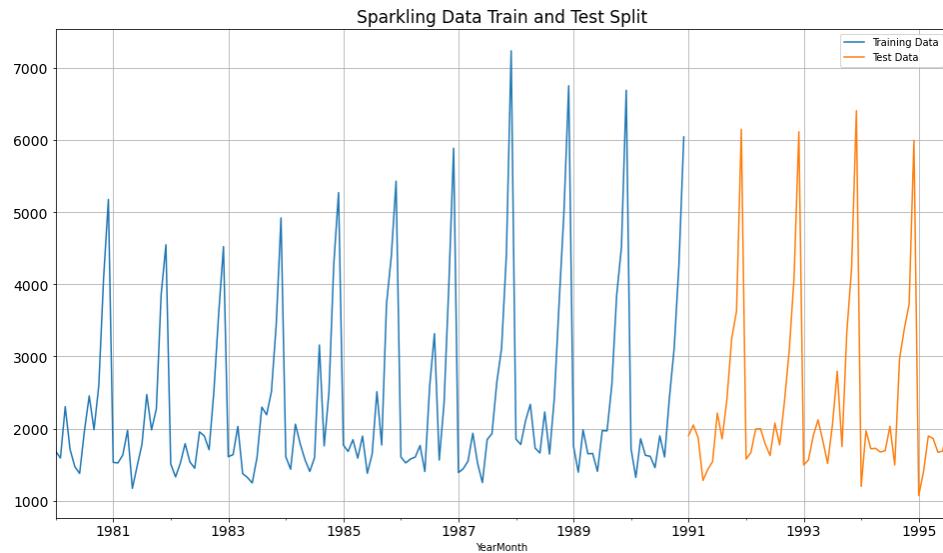
Sparkling	Sparkling	
YearMonth	YearMonth	
1980-01-01	1990-08-01	1686
1980-02-01	1990-09-01	1591
1980-03-01	1990-10-01	2304
1980-04-01	1990-11-01	1712
1980-05-01	1990-12-01	1471
First few rows of Sparkling Training Data		
1990-08-01	1605	
1990-09-01	2424	
1990-10-01	3116	
1990-11-01	4286	
1990-12-01	6047	

Last few rows of Sparkling Test Data

Sparkling	
YearMonth	
1995-03-01	1897
1995-04-01	1862
1995-05-01	1670
1995-06-01	1688
1995-07-01	2031

First few rows of Sparkling Test Data

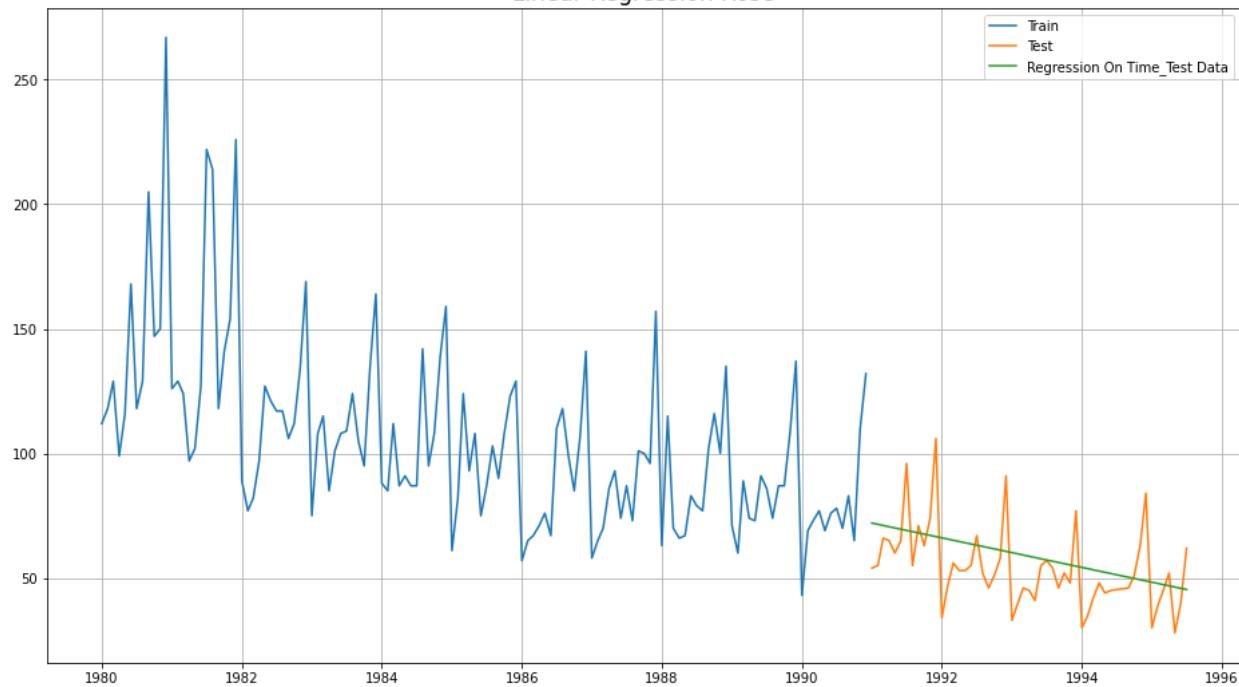
Sparkling	
YearMonth	
1991-01-01	1902
1991-02-01	2049
1991-03-01	1874
1991-04-01	1279
1991-05-01	1432



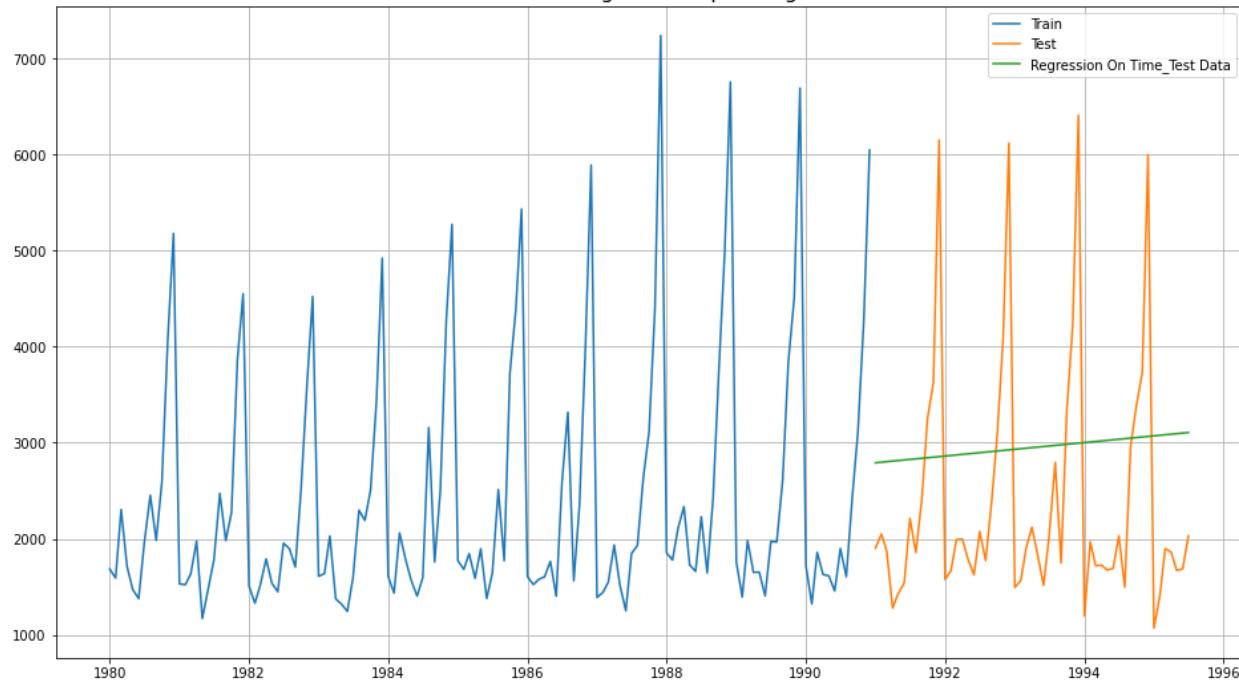
Q 4. Build various exponential smoothing models on the training data and evaluate the model using RMSE on the test data. Other mode Is such as regression, naïve forecast models, simple average models etc. should also be built on the training data and check the performance on the test data using RMSE.

Model 1 - Linear Regression

Linear Regression Rose



Linear Regression Sparkling



Test RMSE Rose Test RMSE Sparkling

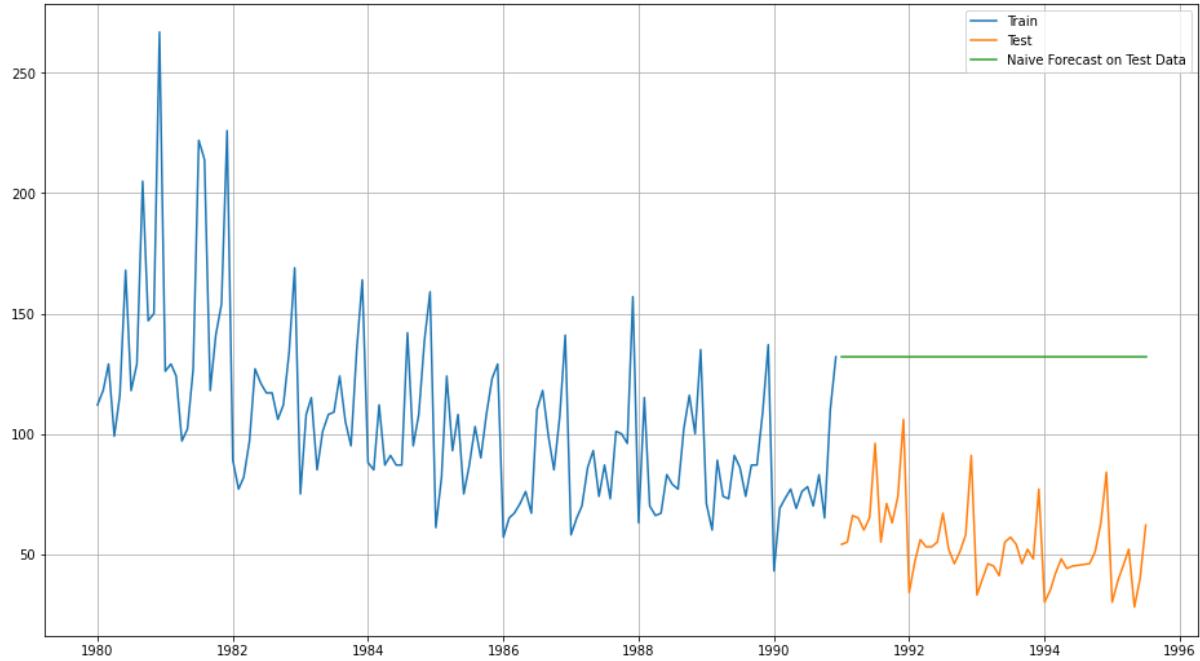
RegressionOnTime

15.268955

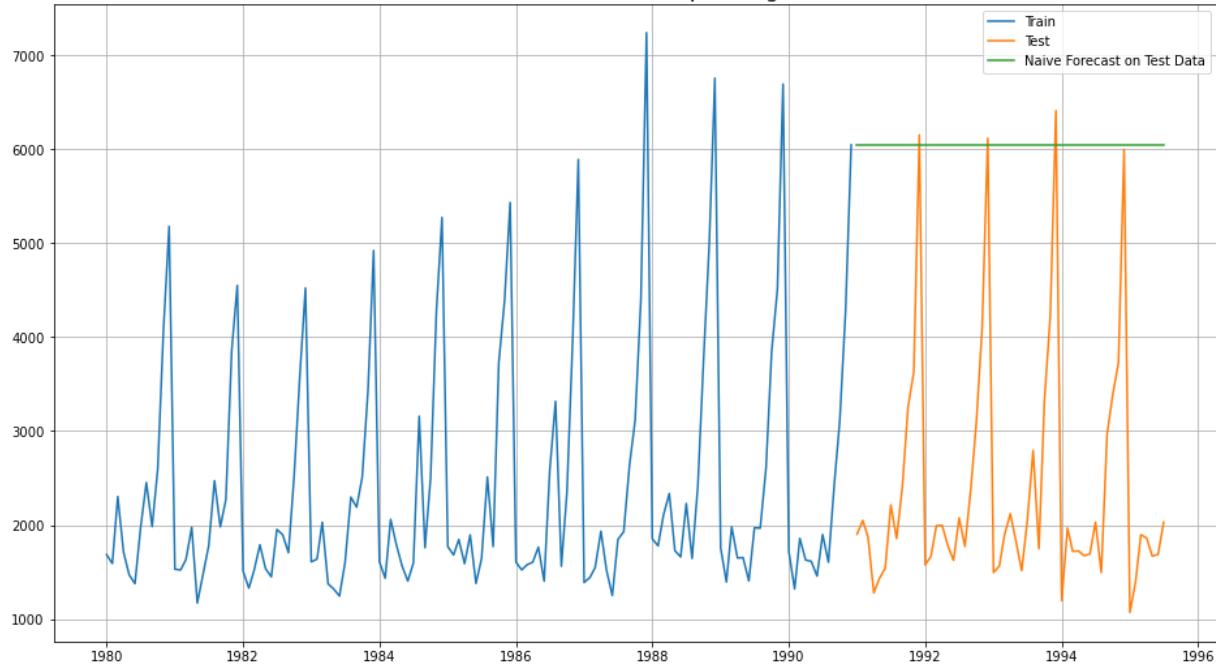
1389.135175

Model 2 - Naive Bayes

Naive Forecast - Rose



Naive Forecast - Sparkling



	Test RMSE Rose	Test RMSE Sparkling
--	----------------	---------------------

NaiveModel	79.718773	3864.279352
------------	-----------	-------------

	Test RMSE Rose	Test RMSE Sparkling
--	----------------	---------------------

RegressionOnTime	15.268955	1389.135175
------------------	-----------	-------------

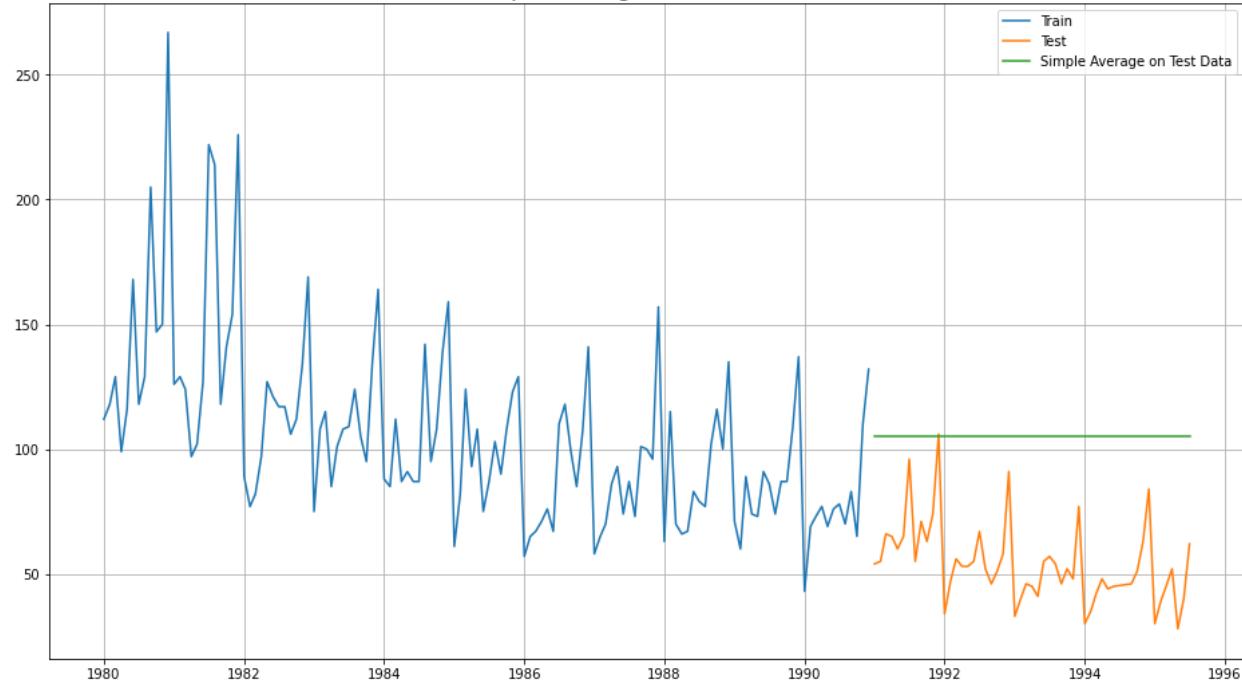
	Test RMSE Rose	Test RMSE Sparkling
--	----------------	---------------------

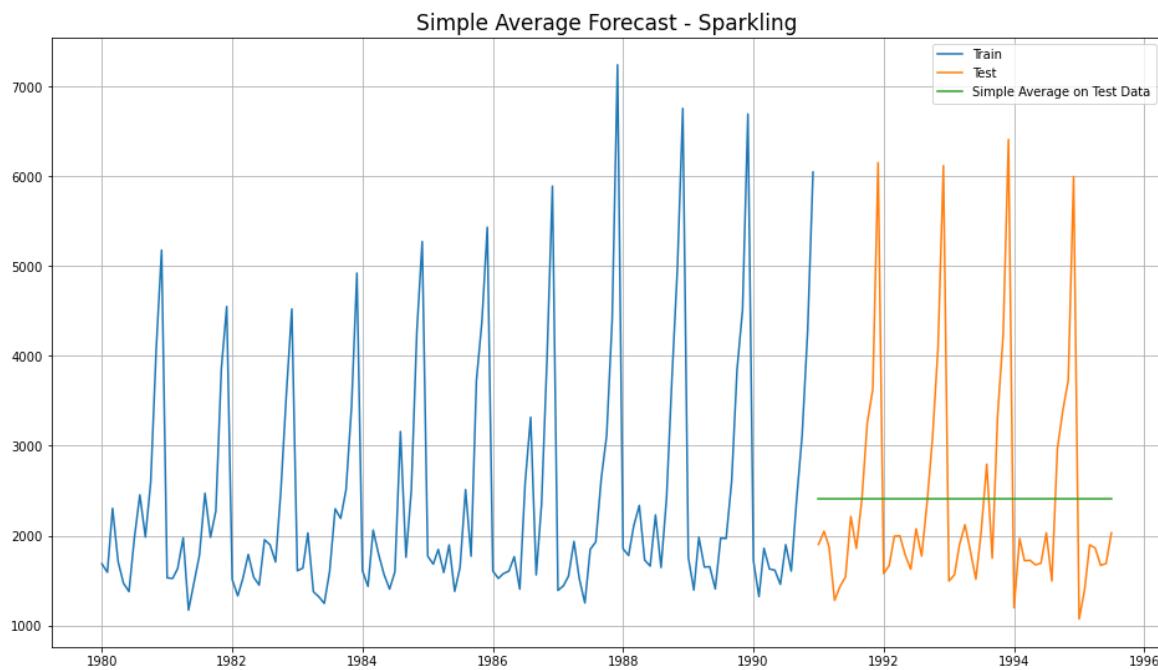
RegressionOnTime	15.268955	1389.135175
------------------	-----------	-------------

NaiveModel	79.718773	3864.279352
------------	-----------	-------------

Model 3 - Simple Average

Simple Average Forecast - Rose

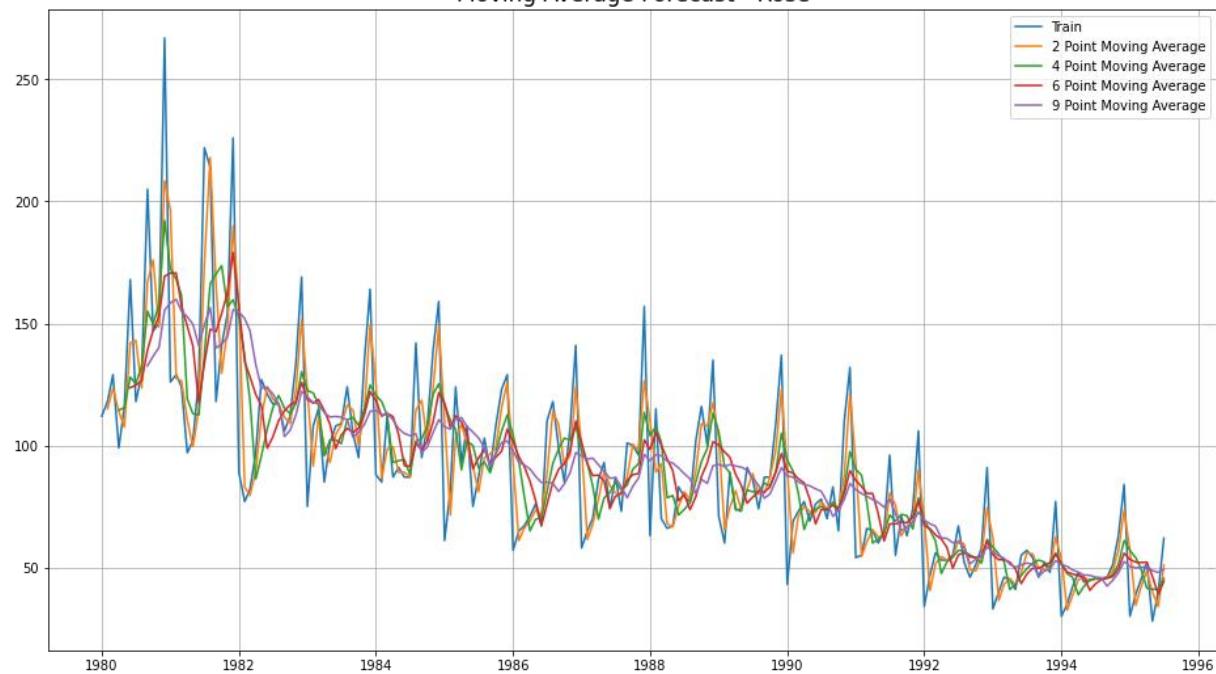




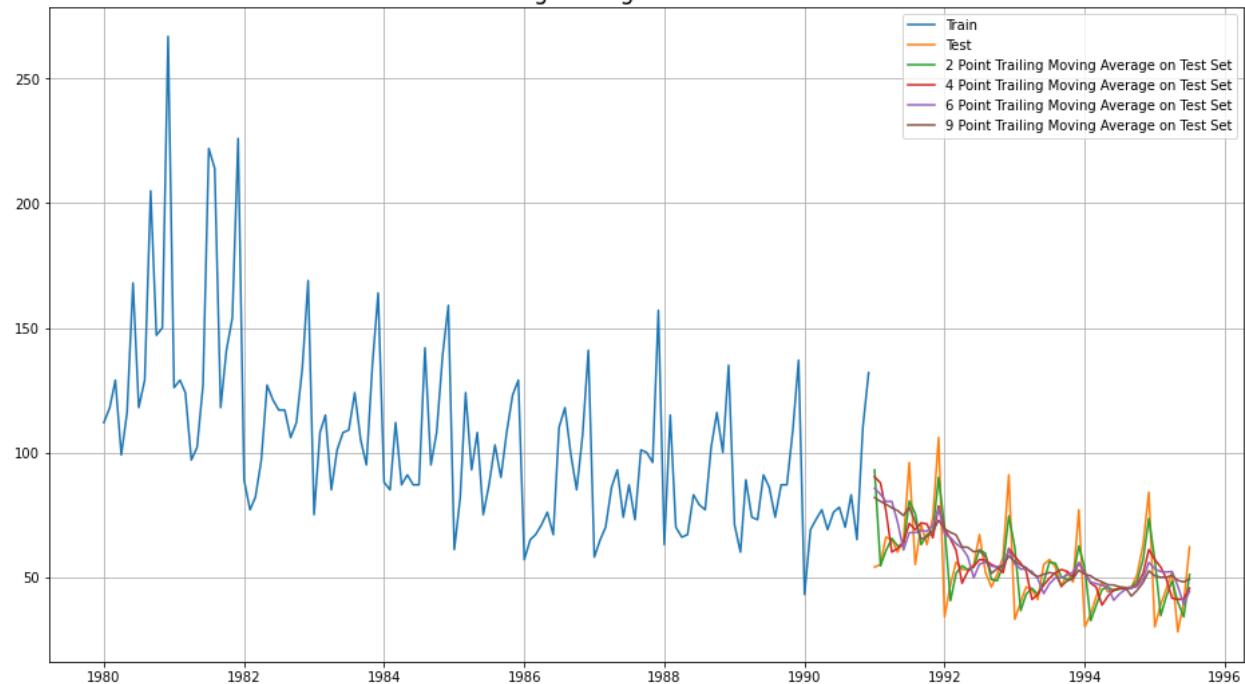
	Test RMSE Rose	Test RMSE Sparkling
RegressionOnTime	15.268955	1389.135175
NaiveModel	79.718773	3864.279352
SimpleAverageModel	53.460570	1275.081804

Model 4.A - Moving Average (Rose)

Moving Average Forecast - Rose



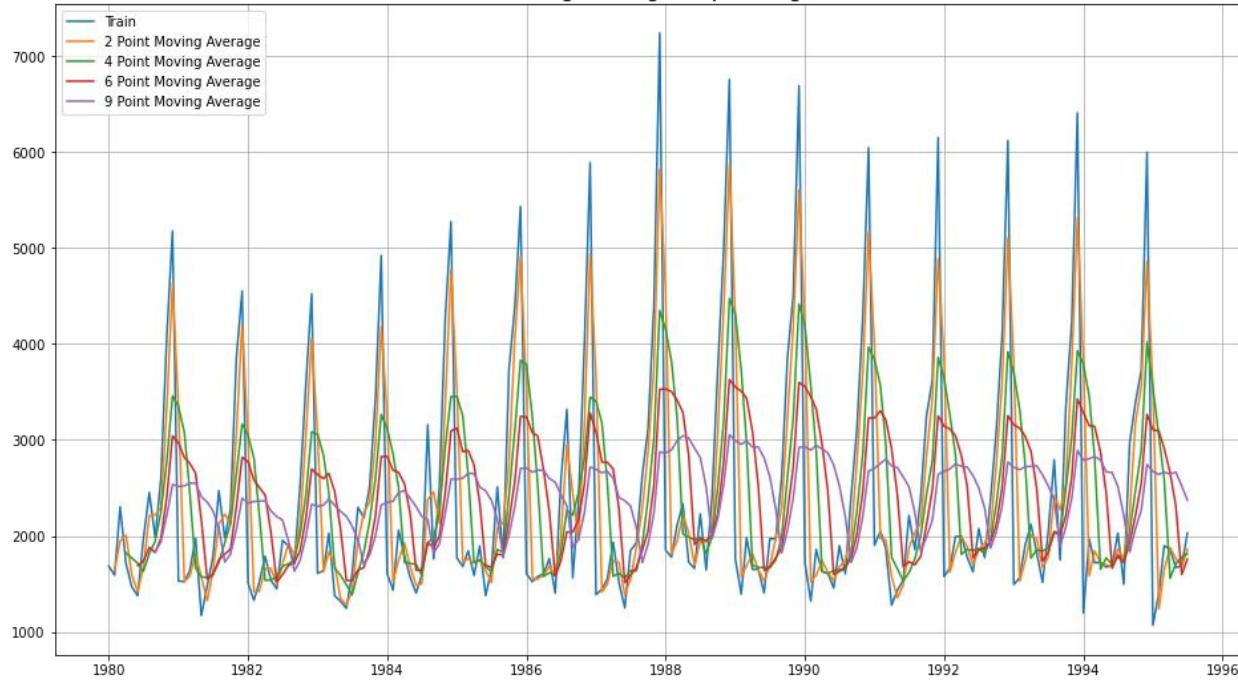
Moving Average Forecast - Rose



Test RMSE Rose

2pointTrailingMovingAverage	11.529278
4pointTrailingMovingAverage	14.451403
6pointTrailingMovingAverage	14.566327
9pointTrailingMovingAverage	14.727630

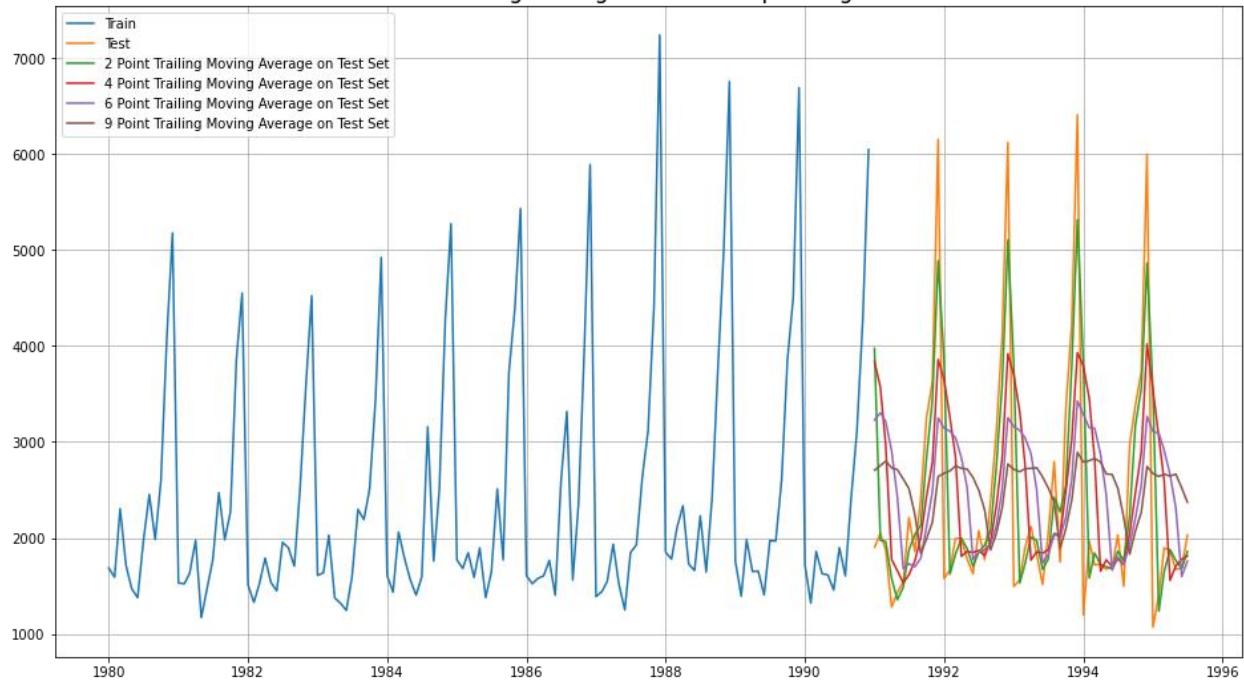
Moving Average - Sparkling



Test RMSE Sparkling

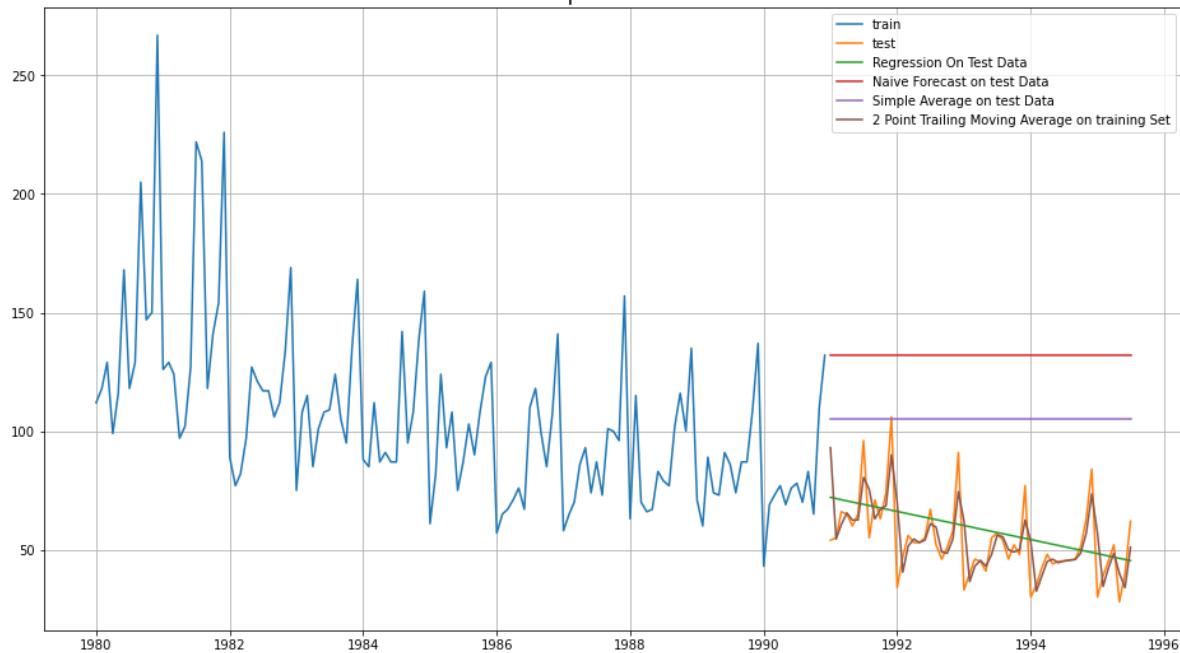
2pointTrailingMovingAverage	813.400684
4pointTrailingMovingAverage	1156.589694
6pointTrailingMovingAverage	1283.927428
9pointTrailingMovingAverage	1346.278315

Moving Average Forecast - Sparkling

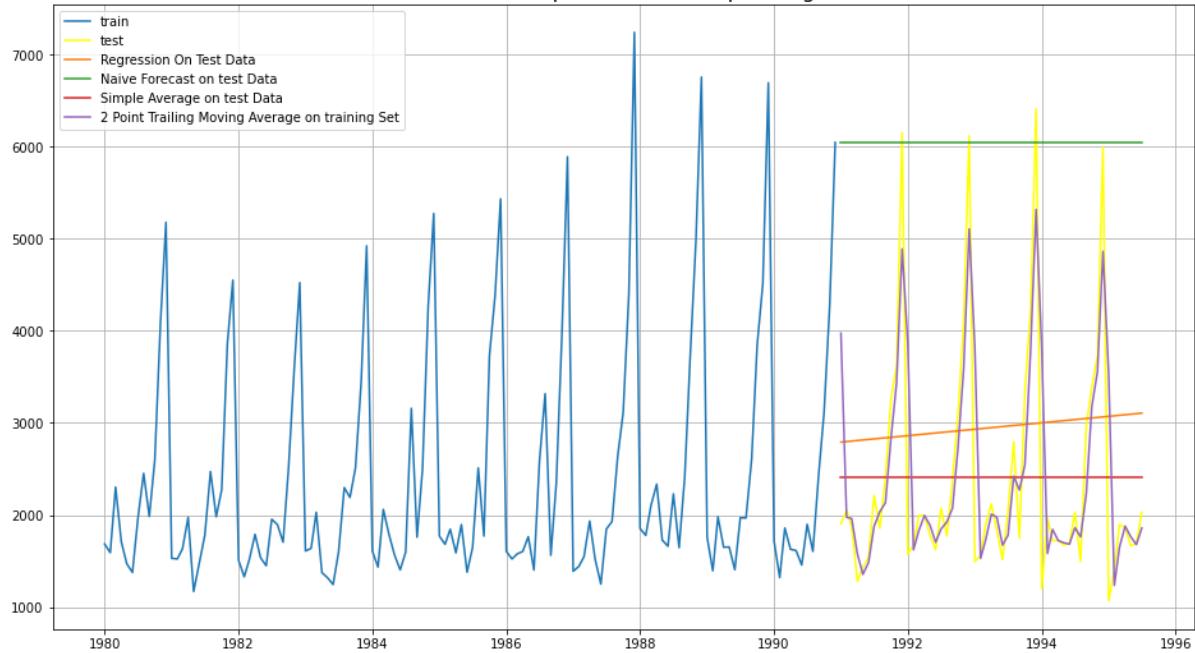


Consolidated Moving Average Forecasts (Rose & Sparkling)

Model Comparison Plots - Rose



Model Comparison Plots - Sparkling



NOTE - • We have built 4 models till now for both Rose and Sparkling Wine datasets

- We fitted various models to the Train split and Tested it on Test split. Accuracy metrics used is Root Mean Squared Error (RMSE) on Test data
- Model 1 - Linear Regression ()

- We regressed variables ‘Rose’ and ‘Sparkling’ against their individual time instances
- We modified the datasets and tagged individual sales to their time instances
- TEST RMSE ROSE = 15.27 | TEST RMSE SPARKLING = 1389.14
- Model 2 - Naive Approach ()
 - Naive approach says that prediction for tomorrow is same as today
 - And, prediction for day-after is same as tomorrow
 - So, effectively all future predictions are going to be same as today
- TEST RMSE ROSE = 79.72 | TEST RMSE SPARKLING = 3864.28
- Model 3 - Simple Average ()
- Model 4 - Moving Average (MA)
 - We calculate rolling means (Moving averages) over different intervals for the whole train data
 - 2 Pt MA ==> means, we find average of 1st and 2nd to predict 3rd similarly, average of 2nd and 3rd to predict 4th and so on

- 4 Pt MA ==> means, we find average of 1st, 2nd, 3rd & 4th to predict 5th also, average of 2nd, 3rd, 4th & 5th to predict 6th and so on $y_t = \beta_0 + \beta_1 X_t + \epsilon_t$ $\hat{y}_{t+1} = y_t$ $\hat{y}_{t+2} = \dots = \hat{y}_{t+n} = \text{Mean}(y_1, y_2, \dots, y_t)$

- 2 PT MA ==> TEST RMSE ROSE = 11.53 | TEST RMSE SPARKLING = 813.40

- 4 PT MA ==> TEST RMSE ROSE = 14.45 | TEST RMSE SPARKLING = 1156.59

- 6 PT MA ==> TEST RMSE ROSE = 14.57 | TEST RMSE SPARKLING = 1283.93

- 9 PT MA ==> TEST RMSE ROSE = 14.73 | TEST RMSE SPARKLING = 1346.28

	Test RMSE Rose	Test RMSE Sparkling
RegressionOnTime	15.268955	1389.135175
NaiveModel	79.718773	3864.279352
SimpleAverageModel	53.460570	1275.081804
2pointTrailingMovingAverage	11.529278	813.400684
4pointTrailingMovingAverage	14.451403	1156.589694
6pointTrailingMovingAverage	14.566327	1283.927428
9pointTrailingMovingAverage	14.727630	1346.278315

Consolidated Scores of Regression, Naive, Simple Average & Moving Average

- Till now, Best Model which gives lowest RMSE score for both Rose and Sparkling is --> 2 Pt Moving Average Model

- We'll continue to forecast using Exponential Smoothing Models for both datasets of Rose and Sparkling Wine Sales

- Exponential smoothing averages or exponentially weighted moving averages consist of forecast based on previous periods data with exponentially declining influence on the older observations

- Exponential smoothing methods consist of special case exponential moving with notation ETS (Error, Trend, Seasonality) where each can be None(N), Additive (N), Additive damped (Ad), Multiplicative (M) or Multiplicative damped (Md)

- One or more parameters control how fast the weights decay. The values of the parameters lie between 0 and 1

- We'll build following Exponential Smoothing Models –

- Single Exponential Smoothing with Additive Errors - ETS(A, N, N)

- Double Exponential Smoothing with Additive Errors, Additive Trends - ETS(A, A, N)

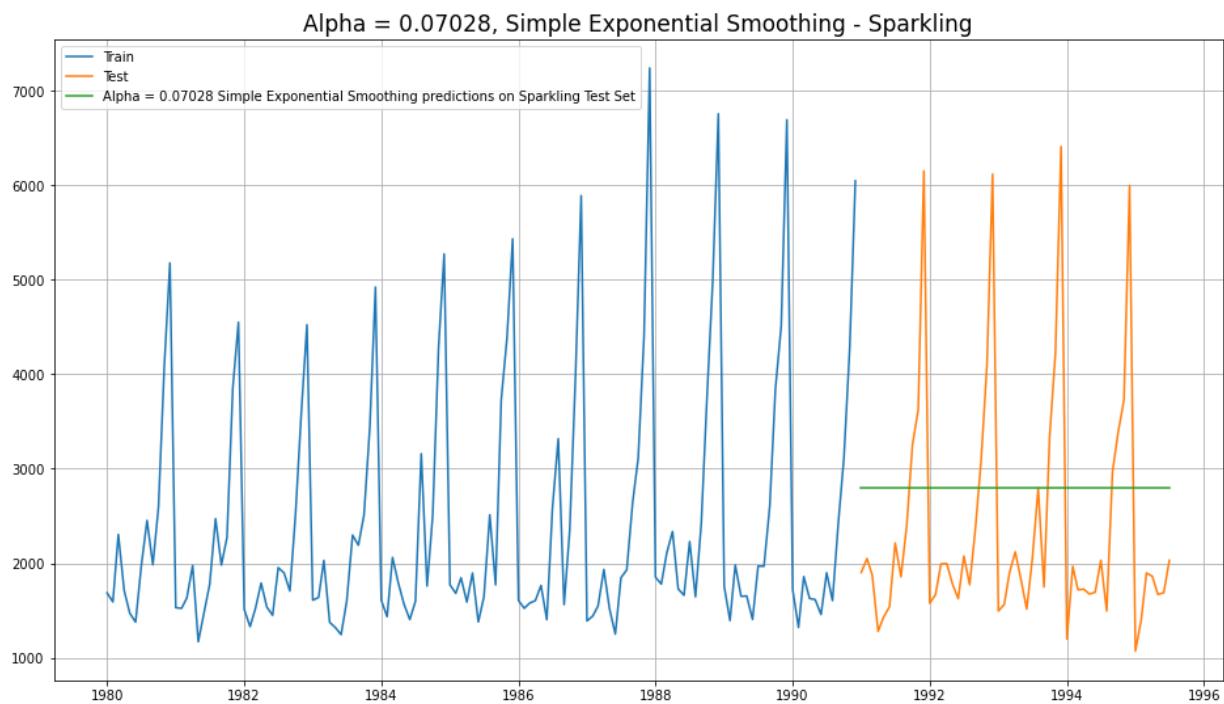
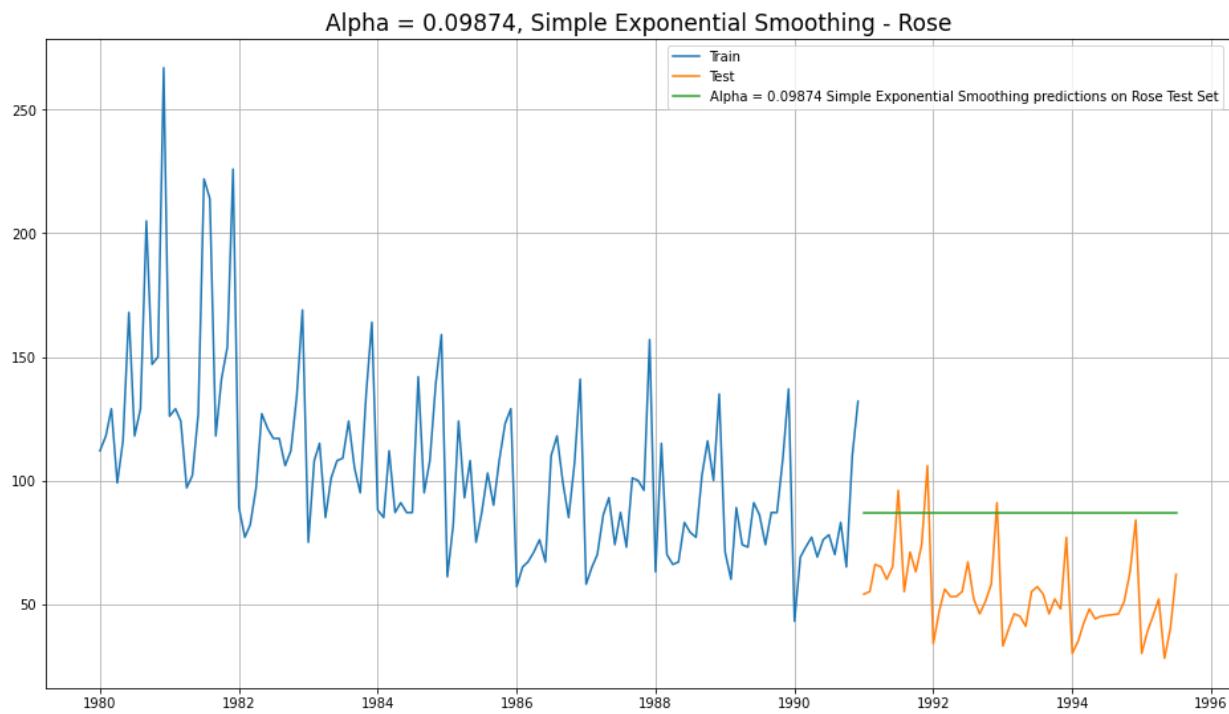
- Triple Exponential Smoothing with Additive Errors, Additive Trends, Additive Seasonality - ETS(A, A, A)

- Triple Exponential Smoothing with Additive Errors, Additive Trends, Multiplicative Seasonality - ETS(A, A, M)

- Triple Exponential Smoothing with Additive Errors, Additive DAMPED Trends, Additive Seasonality - ETS(A, Ad, A)

- Triple Exponential Smoothing with Additive Errors, Additive DAMPED Trends, Multiplicative Seasonality - ETS(A, Ad, M)

Single Exponential Smoothing with Additive Errors - ETS(A, N, N)



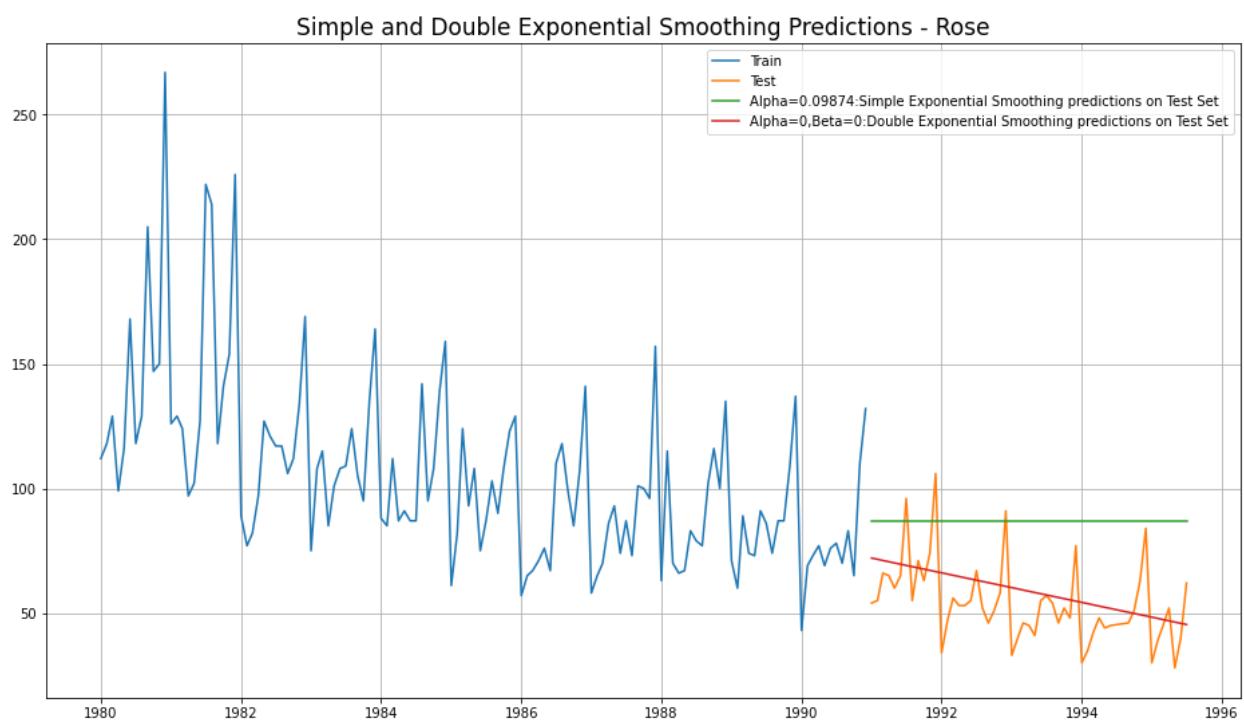
SES RMSE: 36.796235605069505

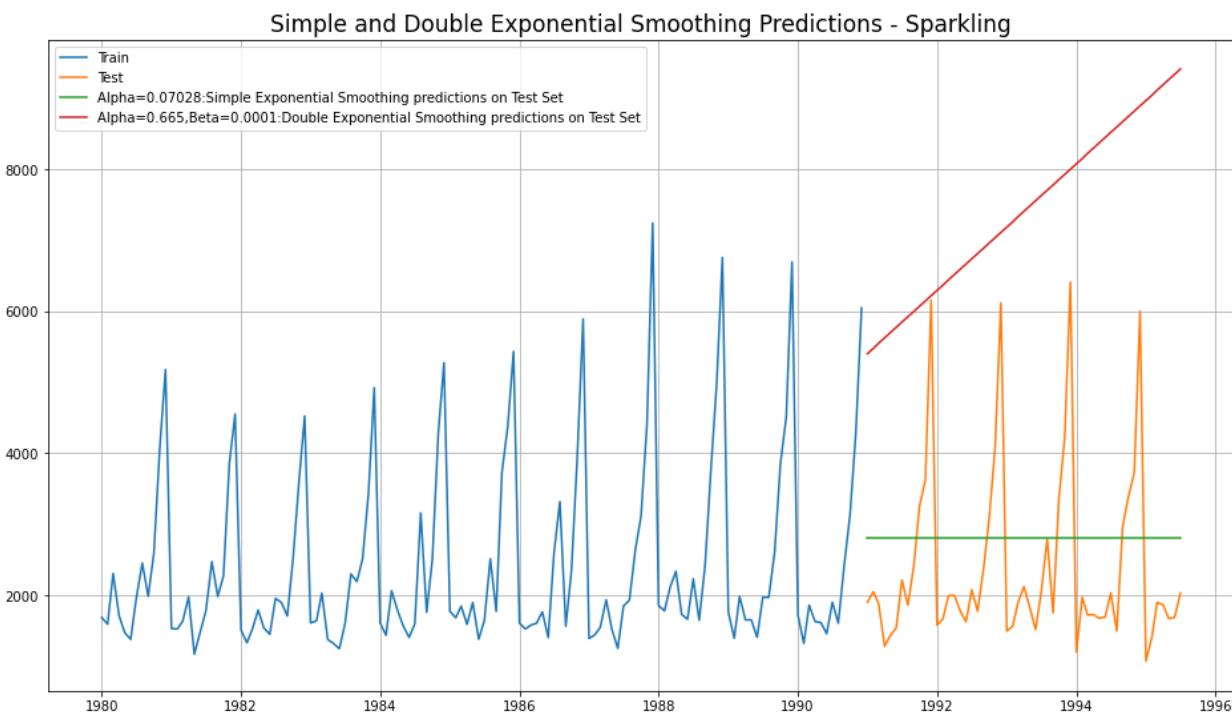
SES RMSE (calculated using statsmodels): 36.7962356050695

	Test RMSE Rose	Test RMSE Sparkling
RegressionOnTime	15.268955	1389.135175
NaiveModel	79.718773	3864.279352
SimpleAverageModel	53.460570	1275.081804
2pointTrailingMovingAverage	11.529278	813.400684
4pointTrailingMovingAverage	14.451403	1156.589694
6pointTrailingMovingAverage	14.566327	1283.927428
9pointTrailingMovingAverage	14.727630	1346.278315
Simple Exponential Smoothing	36.796236	1338.000861

- Best Model till now for Rose and Sparkling ——> 2 Pt Moving Average Model

Double Exponential Smoothing with Additive Errors, Additive Trends - ETS(A, A, N)



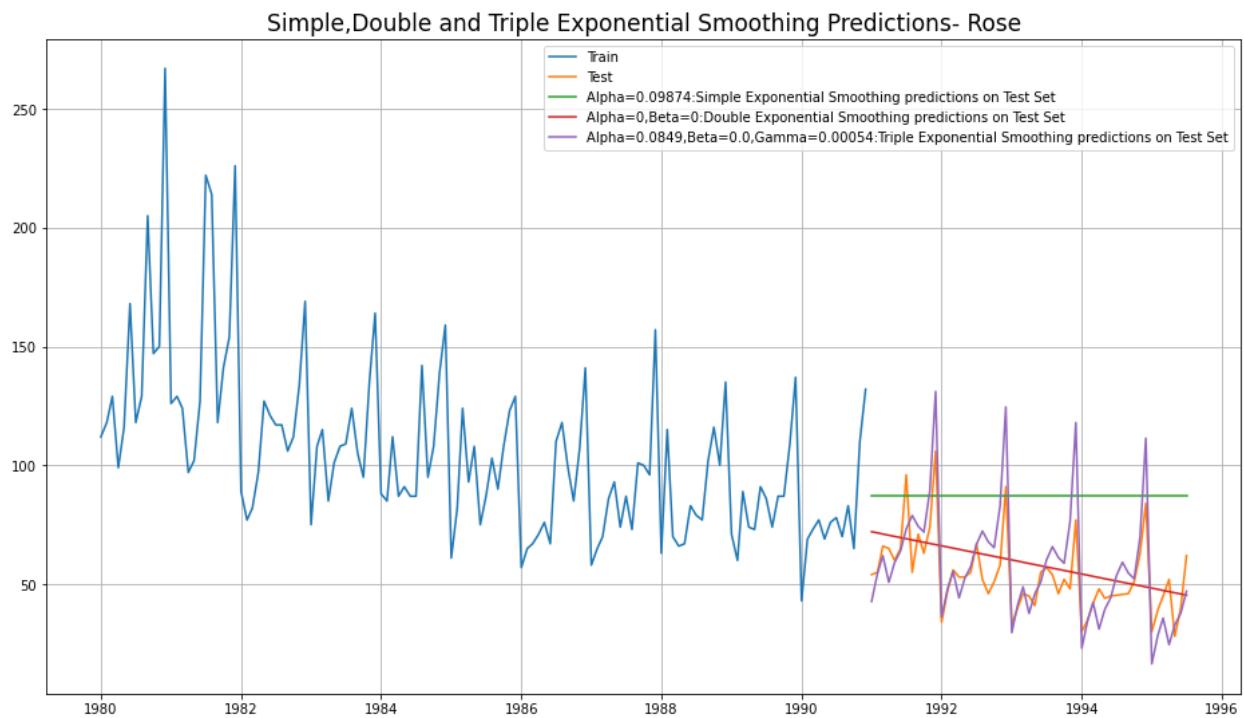


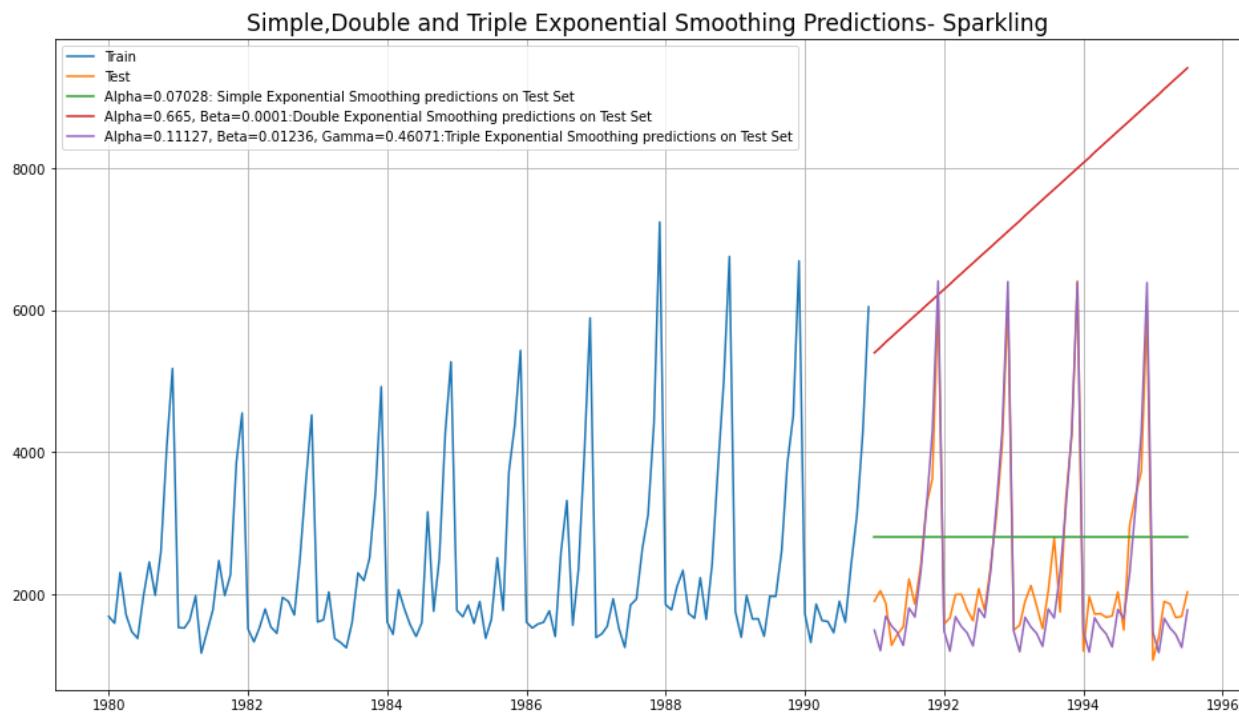
	Test RMSE Rose	Test RMSE Sparkling
RegressionOnTime	15.268955	1389.135175
NaiveModel	79.718773	3864.279352
SimpleAverageModel	53.460570	1275.081804
2pointTrailingMovingAverage	11.529278	813.400684
4pointTrailingMovingAverage	14.451403	1156.589694
6pointTrailingMovingAverage	14.566327	1283.927428
9pointTrailingMovingAverage	14.727630	1346.278315
Simple Exponential Smoothing	36.796236	1338.000861
Double Exponential Smoothing	15.268957	5291.879833

Here, we see that the Double Exponential Smoothing has actually done well when compared to the Simple Exponential Smoothing. This is because of the fact that the Double Exponential Smoothing model has picked up the trend component as well.

- In Rose - DES has picked up the trend well. DES seems to perform better than SES here
- In Sparkling - DES shows a non-existent trend. DES does not perform well here
- Rose - Level parameter, Alpha = 0 Trend parameter, Beta = 0
- Sparkling - Level parameter, Alpha = 0.665 Trend parameter, Beta = 0.0001
- Best Model till now for Rose and Sparkling ——> 2 Pt Moving Average Model

Triple Exponential Smoothing with Additive Errors, Additive Trends, Additive Seasonality - ETS(A, A, A)





	Test RMSE Rose	Test RMSE Sparkling
RegressionOnTime	15.268955	1389.135175
NaiveModel	79.718773	3864.279352
SimpleAverageModel	53.460570	1275.081804
2pointTrailingMovingAverage	11.529278	813.400684
4pointTrailingMovingAverage	14.451403	1156.589694
6pointTrailingMovingAverage	14.566327	1283.927428
9pointTrailingMovingAverage	14.727630	1346.278315
Simple Exponential Smoothing	36.796236	1338.000861
Double Exponential Smoothing	15.268957	5291.879833
Triple Exponential Smoothing (Additive Season)	14.278440	378.625883

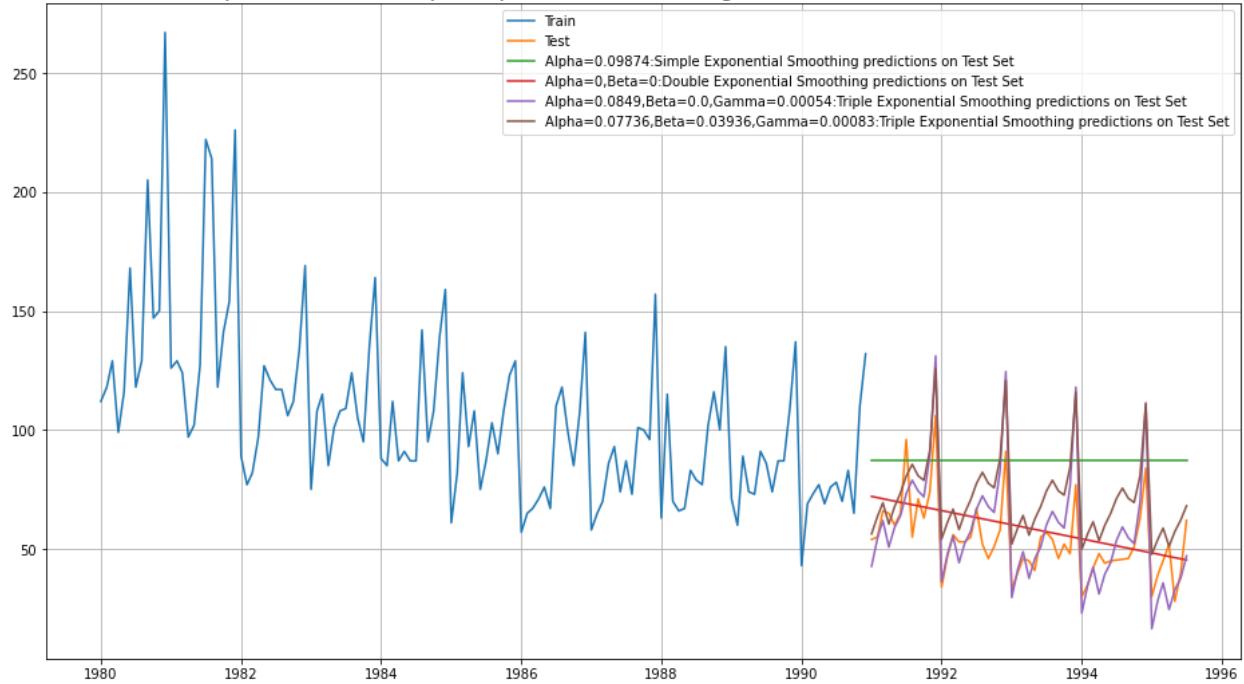
- In Rose & Sparkling - TES has picked up the trend and seasonality very well

- Rose - Level parameter, Alpha = 0.0849 Trend parameter, Beta = 0.0 Seasonality parameter, Gamma = 0.00054

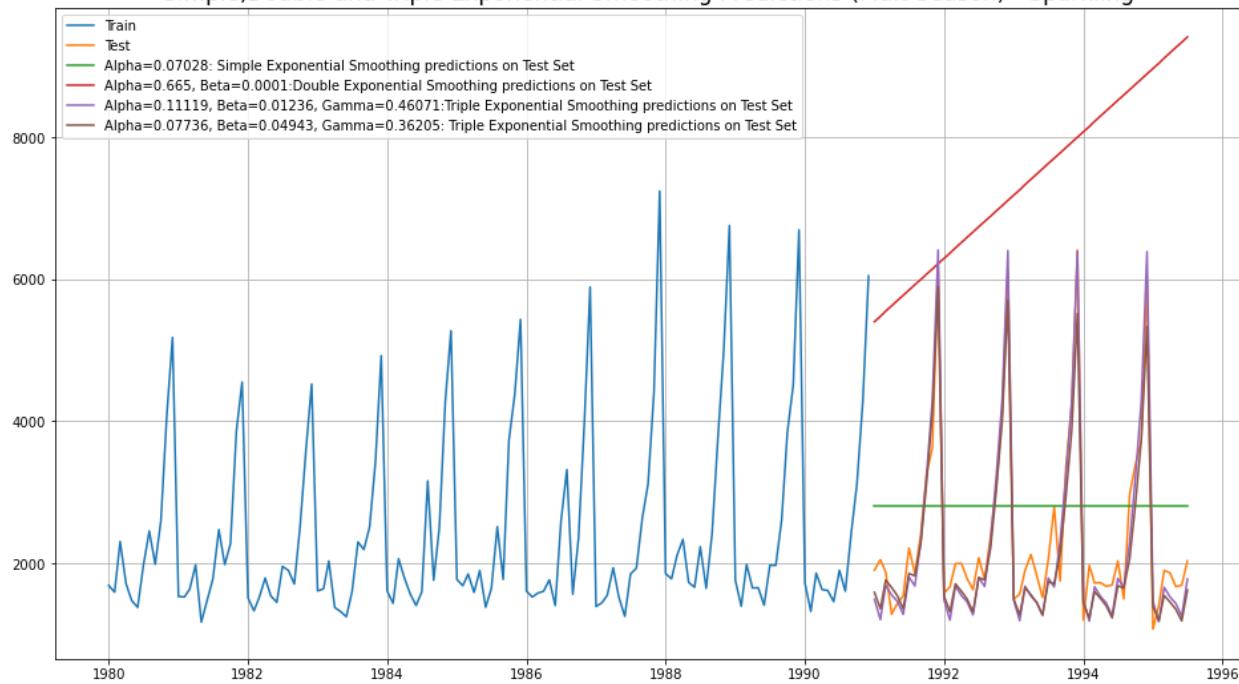
- Sparkling - Level parameter, Alpha = 0.11127 Trend parameter, Beta = 0.01236 Seasonality parameter, Gamma = 0.46071

Triple Exponential Smoothing with Additive Errors, Additive Trends, Multiplicative Seasonality - ETS(A, A, M).

Simple, Double and Triple Exponential Smoothing Predictions (Mult Season) - Rose



Simple,Double and Triple Exponential Smoothing Predictions (Mult Season) - Sparkling



- Rose - Level parameter, Alpha = 0.07736 Trend parameter, Beta = 0.03936 Seasonality parameter, Gamma = 0.00083
- Sparkling - Level parameter, Alpha = 0.07736 Trend parameter, Beta = 0.04943 Seasonality parameter, Gamma = 0.36205

		Test RMSE Rose	Test RMSE Sparkling
	RegressionOnTime	15.268955	1389.135175
	NaiveModel	79.718773	3864.279352
	SimpleAverageModel	53.460570	1275.081804
	2pointTrailingMovingAverage	11.529278	813.400684
	4pointTrailingMovingAverage	14.451403	1156.589694
	6pointTrailingMovingAverage	14.566327	1283.927428
	9pointTrailingMovingAverage	14.727630	1346.278315
	Simple Exponential Smoothing	36.796236	1338.000861
	Double Exponential Smoothing	15.268957	5291.879833
	Triple Exponential Smoothing (Additive Season)	14.278440	378.625883
	Triple Exponential Smoothing (Multiplicative Season)	20.189764	402.936179

Till now, Best Model for Rose —> 2 Pt Moving Average Best Model for Sparkling —>

Holt-Winter - ETS (A, A, A)

Triple Exponential Smoothing with Additive Errors, Additive DAMPED Trends, Additive Seasonality - ETS(A, Ad, A)

- Rose - Level parameter, Alpha = 0.07842

Trend parameter, Beta = 0.01153

Seasonality parameter, Gamma = 0.07738

Damping factor = 0.97503

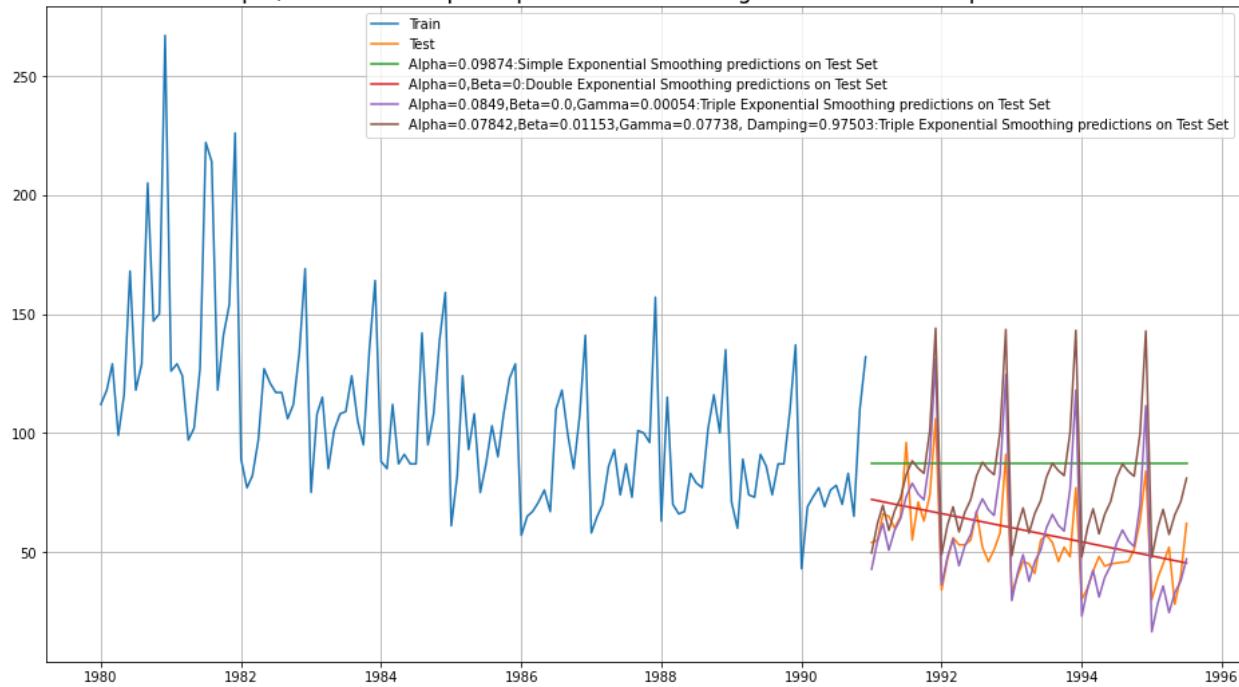
Sparkling - Level parameter, Alpha = 0.10062

Trend parameter, Beta = 0.00018

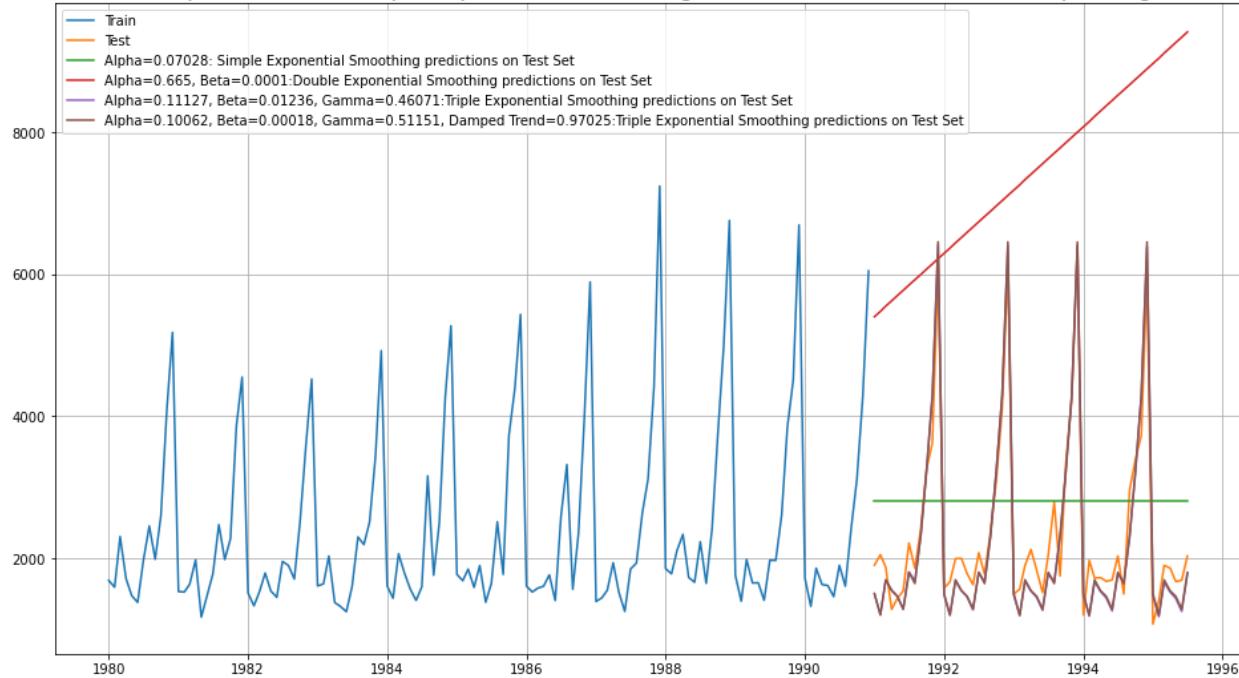
Seasonality parameter, Gamma = 0.97025

Damping factor = 0.97025

Simple,Double and Triple Exponential Smoothing Predictions & Damped - Rose



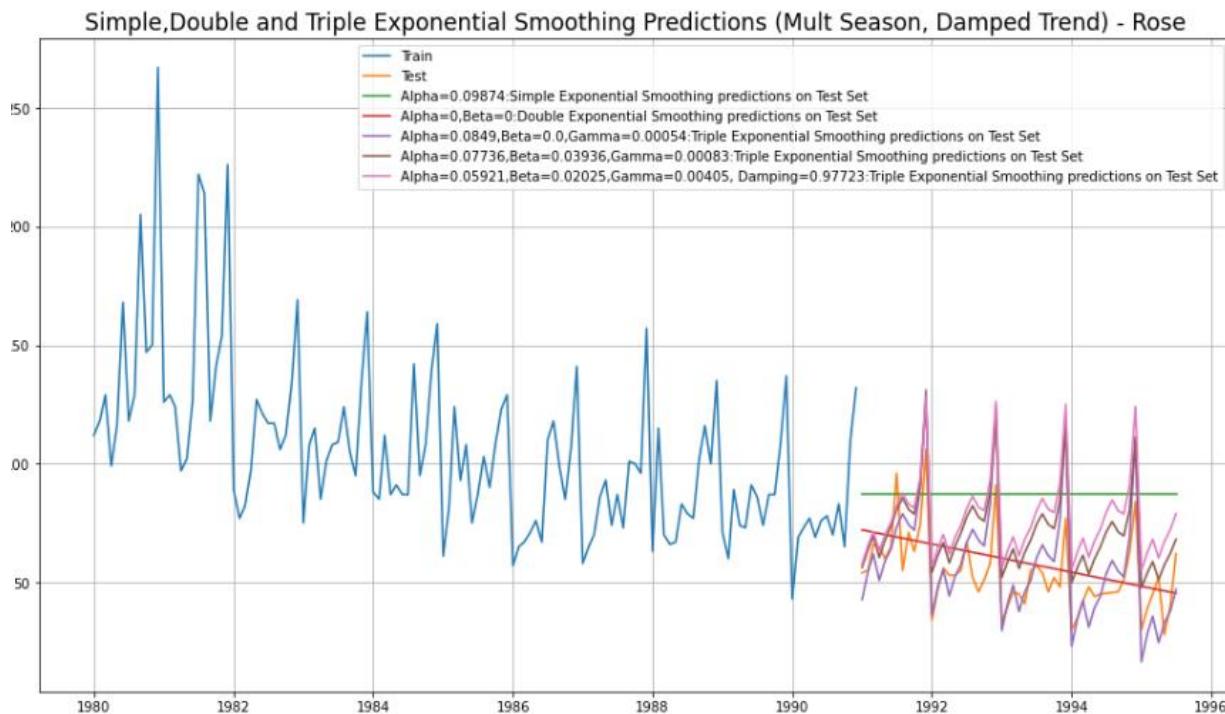
Simple,Double and Triple Exponential Smoothing Predictions (DAMPED TREND)- Sparkling



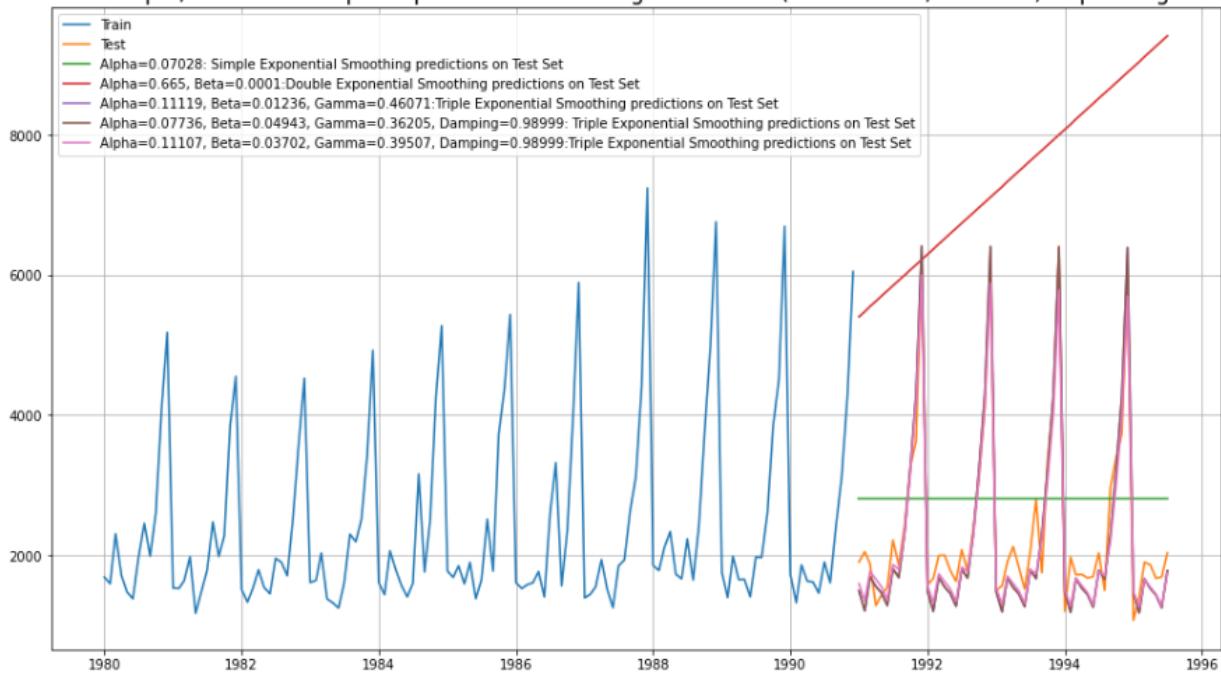
Till now, Best Model for Rose —> 2 Pt Moving Average Best Model for Sparkling —> Holt-Winter - ETS (A, A, A)

	Test RMSE Rose	Test RMSE Sparkling
RegressionOnTime	15.268955	1389.135175
NaiveModel	79.718773	3864.279352
SimpleAverageModel	53.460570	1275.081804
2pointTrailingMovingAverage	11.529278	813.400684
4pointTrailingMovingAverage	14.451403	1156.589694
6pointTrailingMovingAverage	14.566327	1283.927428
9pointTrailingMovingAverage	14.727630	1346.278315
Simple Exponential Smoothing	36.796236	1338.000861
Double Exponential Smoothing	15.268957	5291.879833
Triple Exponential Smoothing (Additive Season)	14.278440	378.625883
Triple Exponential Smoothing (Multiplicative Season)	20.189764	402.936179
Triple Exponential Smoothing (Additive Season, Damped Trend)	28.443719	378.625883

Triple Exponential Smoothing with Additive Errors, Additive DAMPED Trends, Multiplicative Seasonality - ETS(A, Ad, M)



Simple,Double and Triple Exponential Smoothing Predictions (Mult Season, DAMPED) - Sparkling



- Rose - Level parameter, Alpha = 0.05921

Trend parameter, Beta = 0.02025

Seasonality parameter, Gamma = 0.00405

Damping factor = 0.97723

- Sparkling - Level parameter, Alpha = 0.11107

Trend parameter, Beta = 0.03702 Seasonality parameter, Gamma = 0.39507

Damping factor = 0.98999

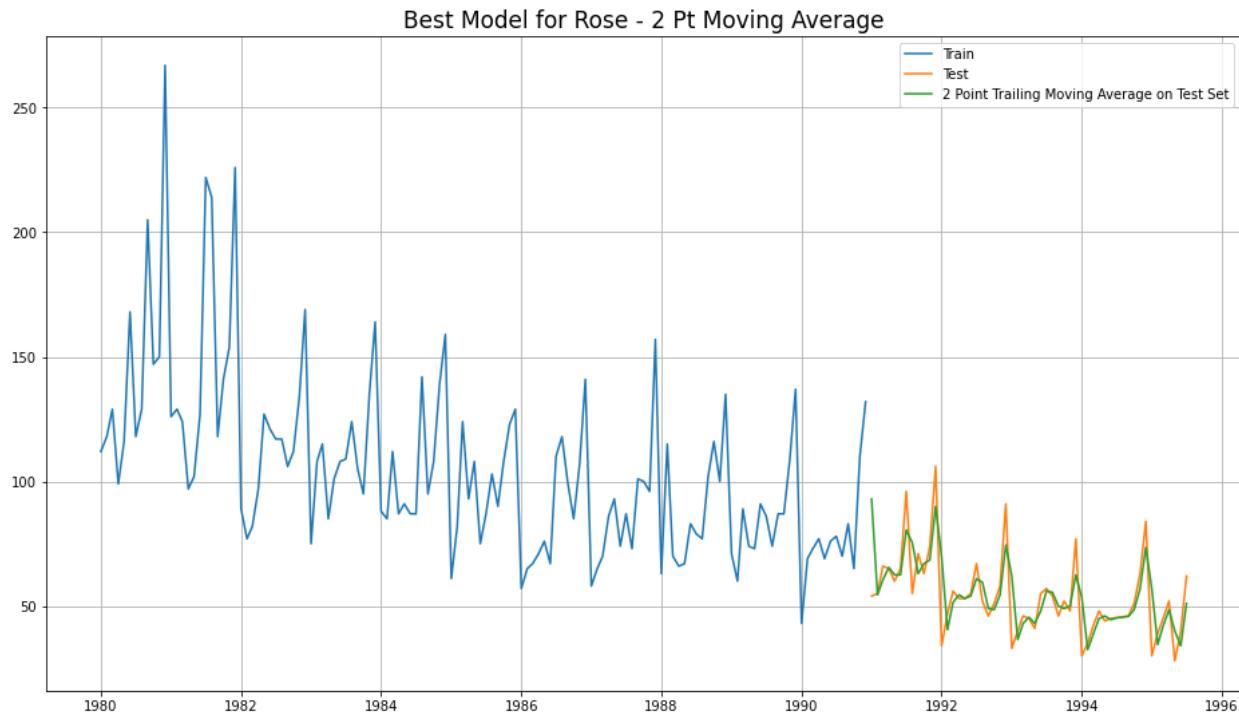
- We conclude that models with least RMSE, Best Model for Rose —> 2 Pt Moving Average
- Best Model for Sparkling —> Holt-Winter Damped Trend ETS (A, Ad, M)

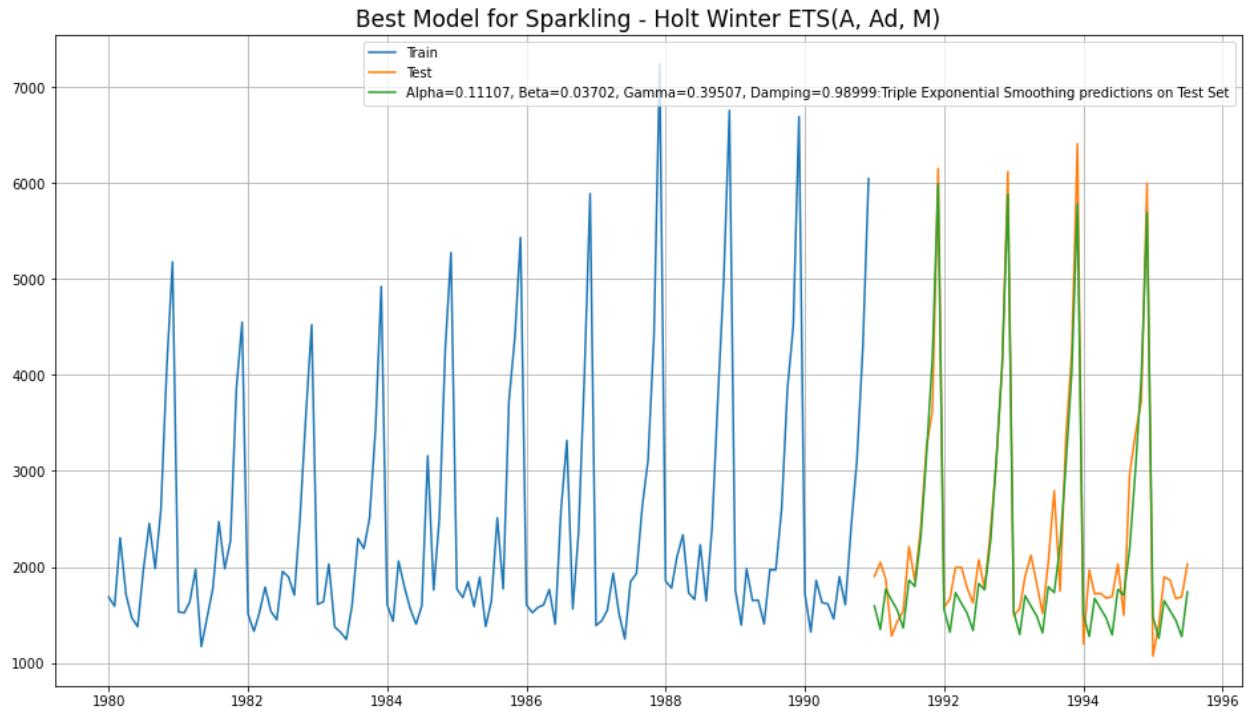
	Test RMSE Rose	Test RMSE Sparkling
RegressionOnTime	15.268955	1389.135175
NaiveModel	79.718773	3864.279352
SimpleAverageModel	53.460570	1275.081804
2pointTrailingMovingAverage	11.529278	813.400684
4pointTrailingMovingAverage	14.451403	1156.589694
6pointTrailingMovingAverage	14.566327	1283.927428
9pointTrailingMovingAverage	14.727630	1346.278315
Simple Exponential Smoothing	36.796236	1338.000861
Double Exponential Smoothing	15.268957	5291.879833
Triple Exponential Smoothing (Additive Season)	14.278440	378.625883
Triple Exponential Smoothing (Multiplicative Season)	20.189764	402.936179
Triple Exponential Smoothing (Additive Season, Damped Trend)	28.443719	378.625883
Triple Exponential Smoothing (Multiplicative Season, Damped Trend)	25.808079	352.443503

We conclude that models with least RMSE, Best Model for Rose —> 2 Pt Moving Average

Best Model for Sparkling —> Holt-Winter Damped Trend ETS (A, Ad, M)

Best Models for Rose and Sparkling-





Q 5] Check for the stationarity of the data on which the model is being built on using appropriate statistical tests and also mention the hypothesis for the statistical test. If the data is found to be non-stationary, take appropriate steps to make it stationary. Check the new data for stationarity and comment. Note: Stationarity should be checked at alpha = 0.05.

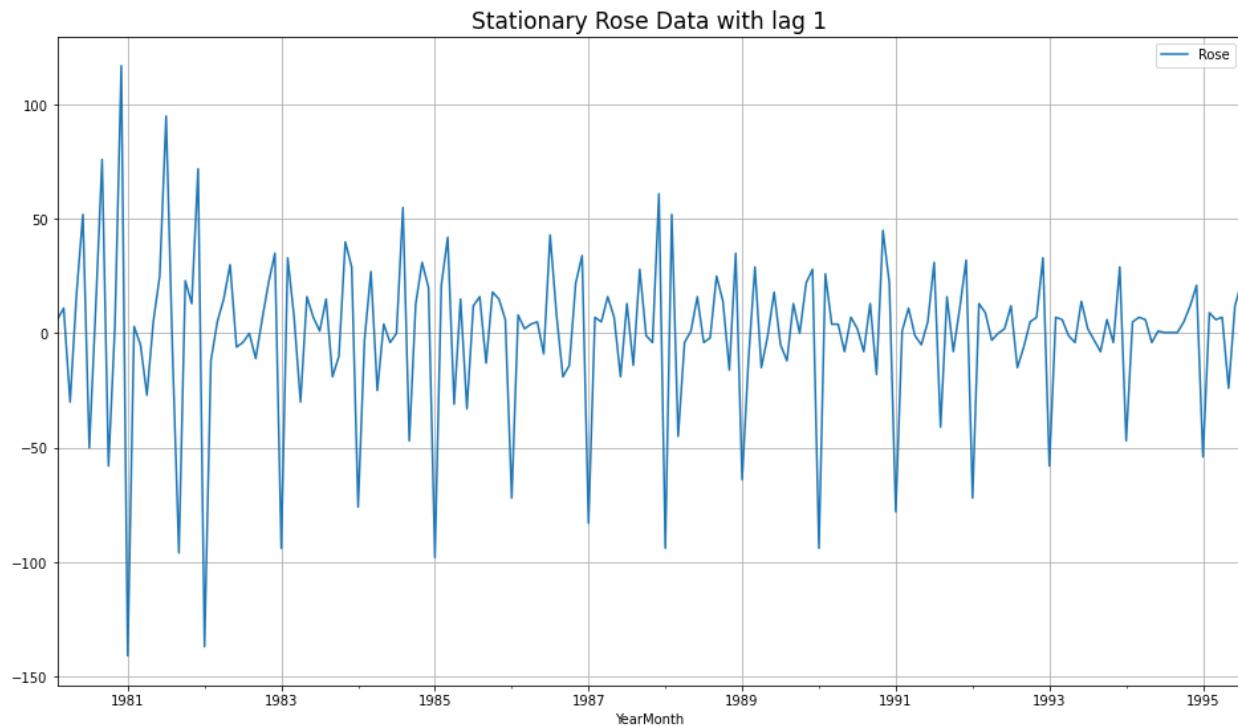
To Check Stationarity of Data –

- We use Augmented Dicky - Fuller (ADF) Test to check the Stationarity of Data
- Hypotheses of ADF Test : Time Series is not Stationary Time Series is Stationary
- So for Industry standard (also given for this problem), the Confidence Interval is 95%
- Hence, alpha = 0.05
- So in ADF Test, if p-value < alpha ==> We reject the Null Hypothesis and hence conclude that given Time Series is Stationary
- So in ADF Test, if p-value > alpha ==> We fail to reject the Null Hypothesis and hence conclude that given Time Series is Not Stationary

- If Time Series is not Stationary then we apply one level of differencing and check for Stationarity again.
- Again, if the Time Series is still not Stationary, we apply one more level of differencing and check for Stationarity again
- Generally, with max 2 levels of differencing, Time Series becomes Stationary
- Once the Time Series is Stationary then we are ready to apply ARIMA / SARIMA models

Stationarity of Rose Wine Dataset –

- Augmented Dicky-Fuller Test was applied to the whole Rose dataset
- We found, p-value = 0.4671
- Here, p-value > alpha=0.05
- We fail to reject the Null Hypothesis and hence conclude that Rose Wine Time Series is Not Stationary
- We take 1 level of differencing and check again for Stationarity
- Now, p-value = 3.0159e-11 0.00
- Now, p-value < alpha=0.05
- Now, we reject the Null Hypothesis and conclude that Rose Time Series is Stationary with a lag of 1.

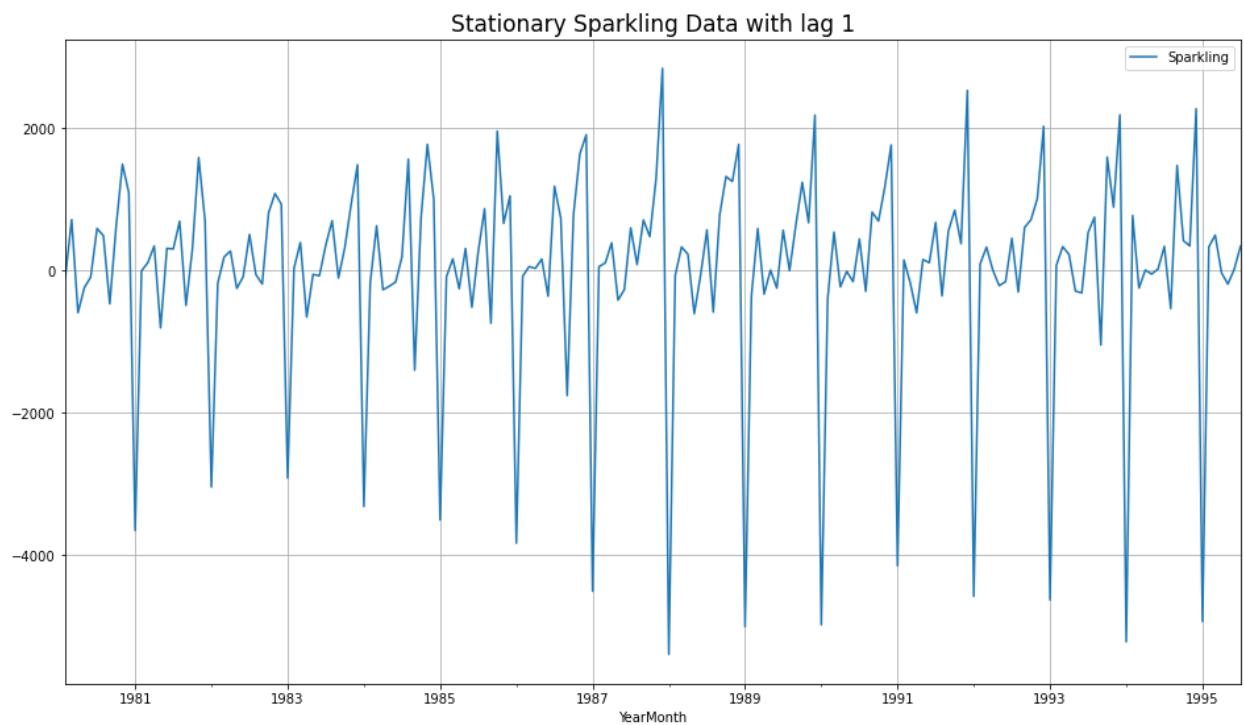


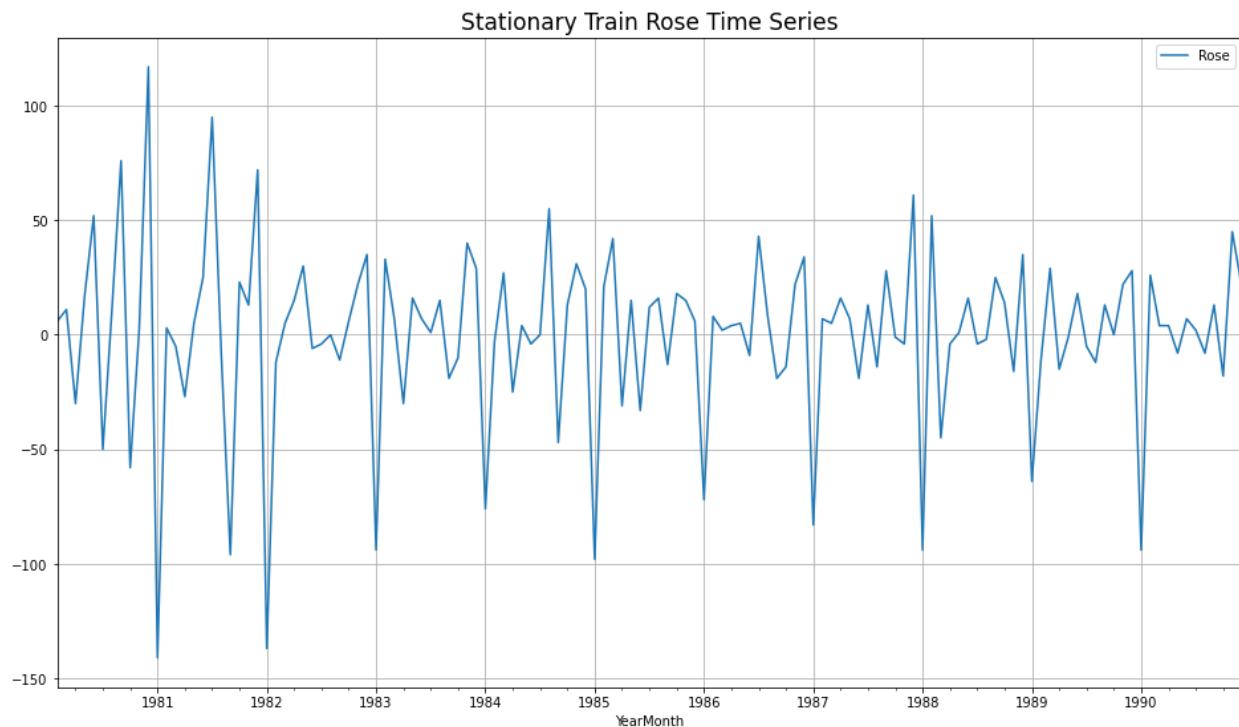
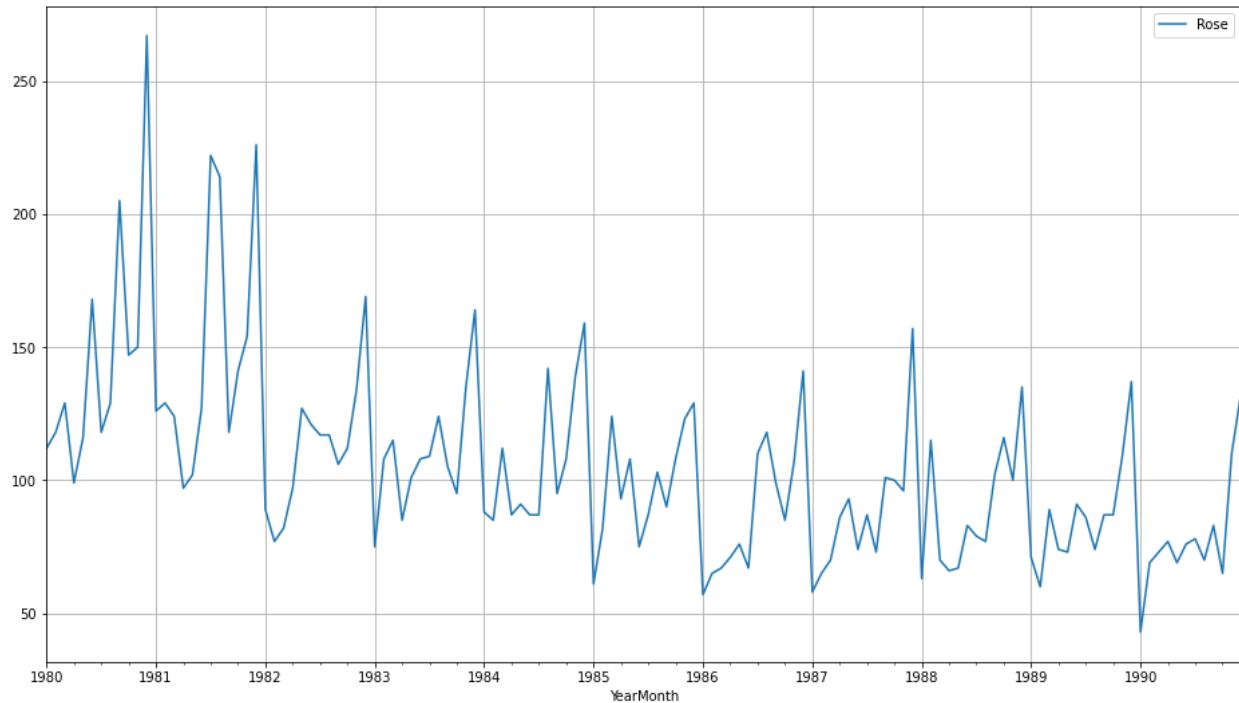
Stationarity of Sparkling Wine Dataset –

- Augmented Dicky-Fuller Test was applied to the whole Sparkling dataset
- We found, p-value = 0.70559
- Here, p-value > alpha=0.05

We fail to reject the Null Hypothesis and hence conclude that Sparkling Wine Time Series is Not Stationary

- We take 1 level of differencing and check again for Stationarity
- Now, p-value = 0.00
- Now, p-value < alpha=0.05
- Now, we reject the Null Hypothesis and conclude that Sparkling Time Series is Stationary with a lag of 1.





Q 6]Build an automated version of the ARIMA/SARIMA model in which the parameters are selected using the lowest Akai

ke Information Criteria (AIC) on the training data and evaluate this model on the test data using RMSE.

ARIMA / SARIMA Models –

- ARIMA is an acronym for Auto-Regressive Integrated Moving Average
- SARIMA stands for Seasonal ARIMA, when the TS has seasonality

ARIMA / SARIMA are forecasting models on Stationary Time Series.

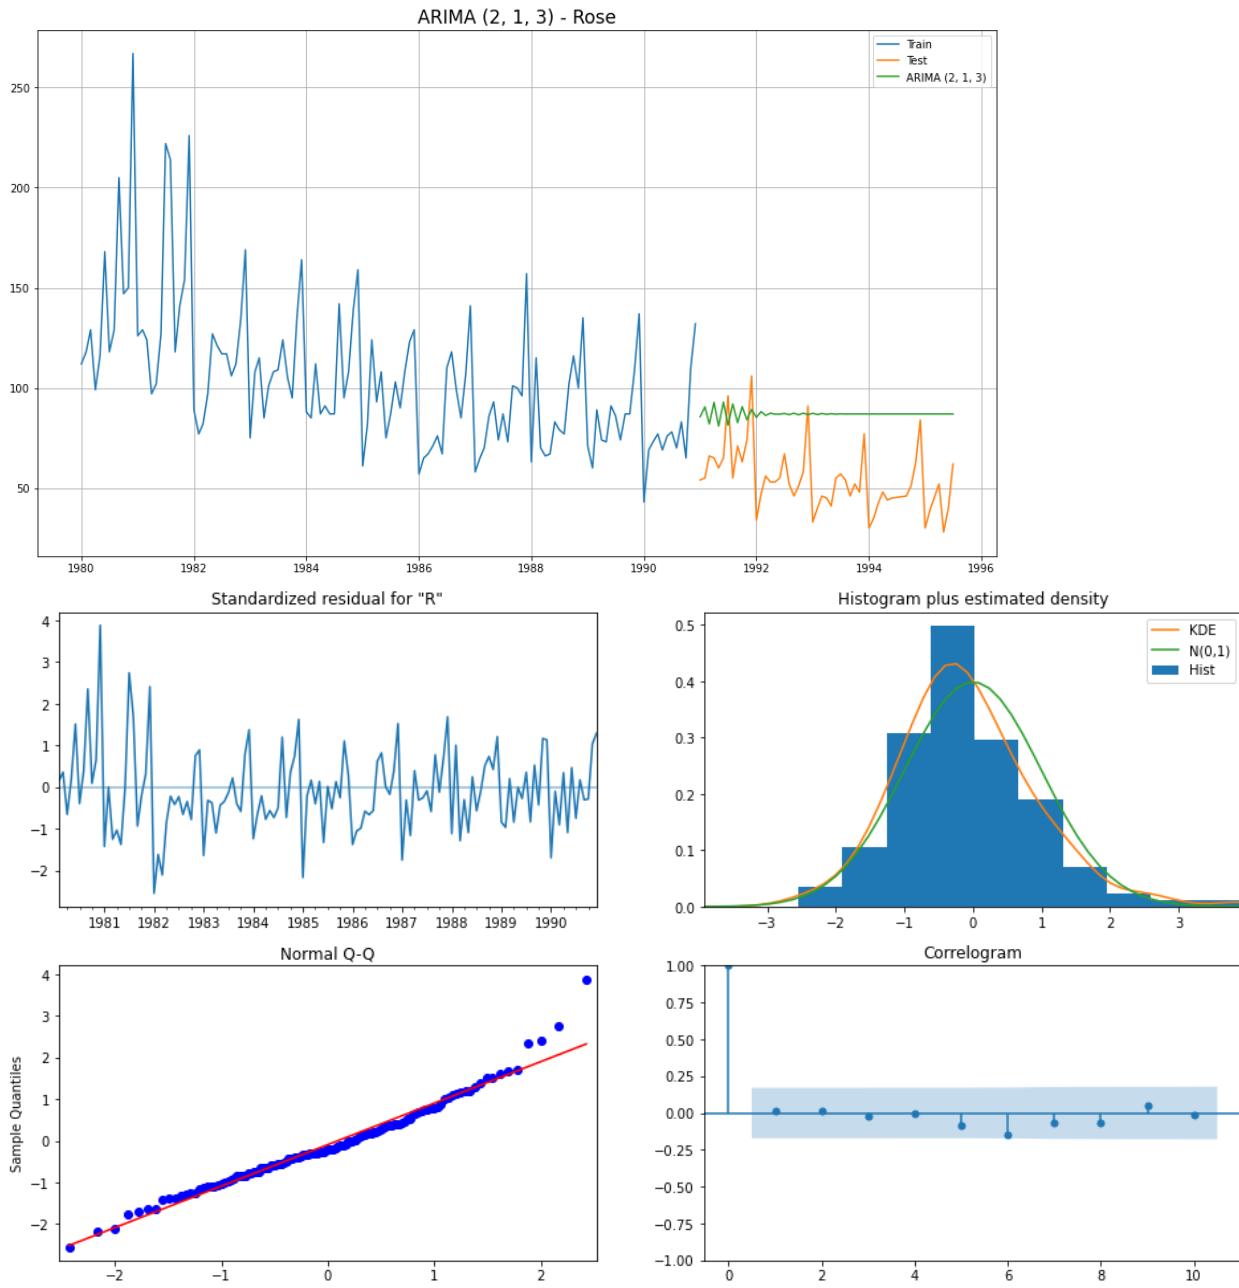
ARIMA / SARIMA Modelling on Train Rose & Sparkling Data –

- We check for stationarity of Train Rose & Sparkling data by using Augmented Dicky Fuller Test
- We take a difference of 1 and make both these datasets Stationary
- We apply the following iterations to both these datasets - 1. ARIMA Automated 2. SARIMA Automated.

ARIMA Automated –

- We create a grid of all possible combinations of (p, d, q)
- Range of p = Range of q = 0 to 3, Constant d = 1
- Few Examples of the grid - Model: (0, 1, 2) Model: (0, 1, 3) Model: (1, 1, 0) Model: (1, 1, 1) Model: (1, 1, 2) Model: (1, 1, 3) Model: (2, 1, 0) Model: (2, 1, 1) Model: (2, 1, 2) Model: (2, 1, 3) Model: (3, 1, 0) Model: (3, 1, 1) • We fit ARIMA models to each of these combinations for both datasets
- We choose the combination with the least Akaike Information Criteria (AIC)
- We fit ARIMA to this combination of (p, d, q) to the Train set and forecast on the Test set
- Finally, we check the accuracy of this model by checking RMSE of Test set.

For Rose, Best Combination with Least AIC is - (p, d, q) —> (2, 1, 3).

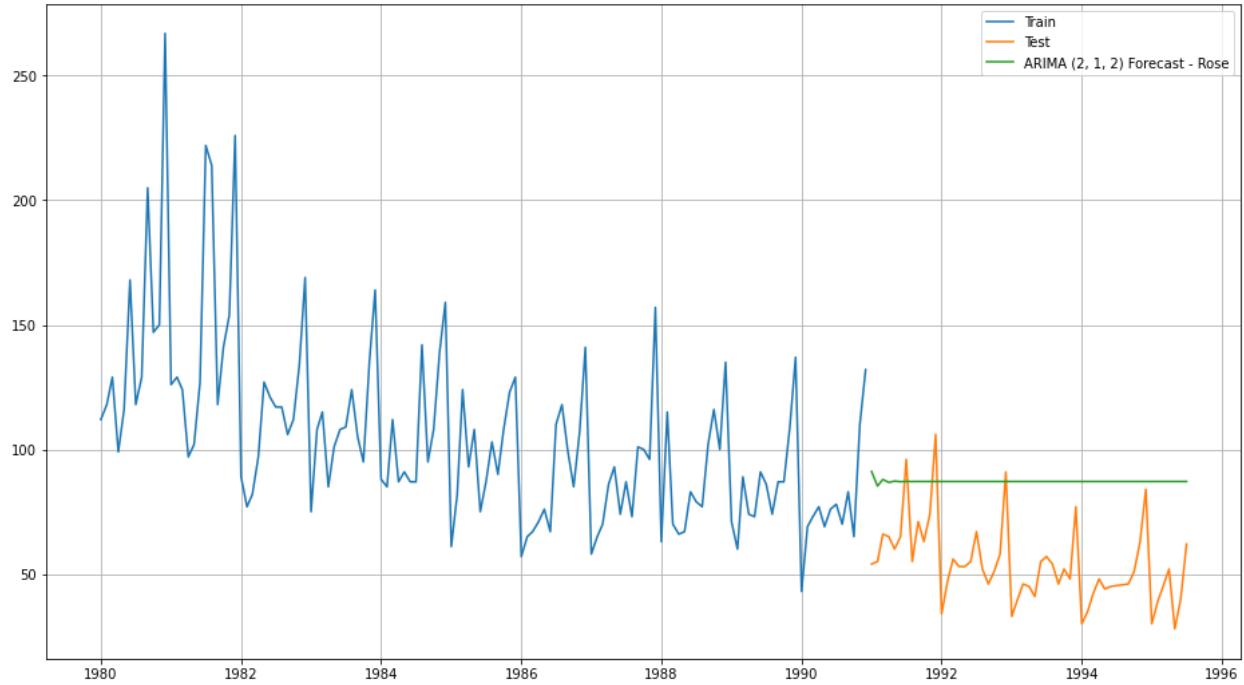


Test RMSE Rose Test MAPE Rose

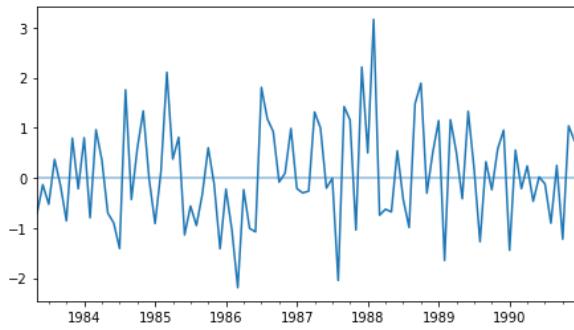
ARIMA(2,1,3)	36.817929	75.849346
---------------------	------------------	------------------

For Sparkling, Best Combination with Least AIC is - (p, d, q) —> (2, 1, 2)

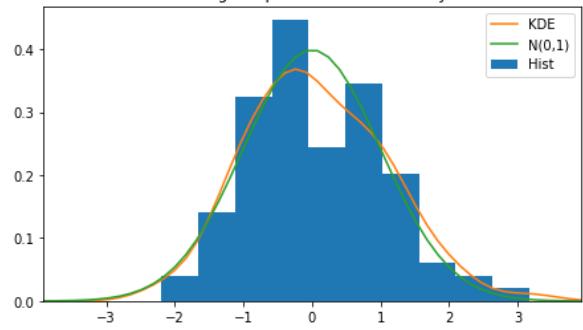
ARIMA (2, 1, 2) Forecast - Rose



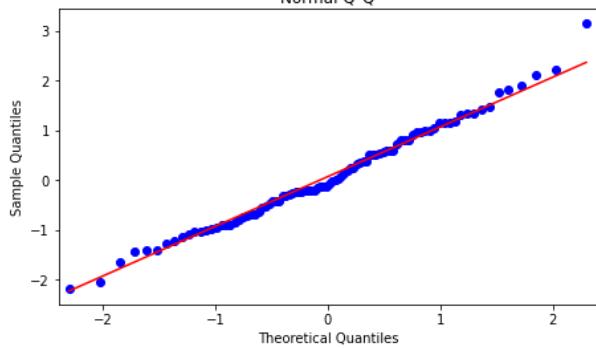
Standardized residual for "R"



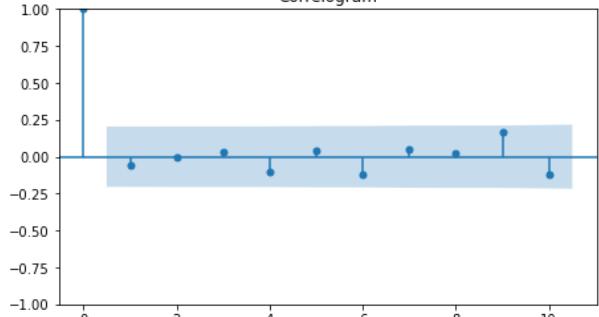
Histogram plus estimated density



Normal Q-Q



Correlogram



Test RMSE Rose	Test MAPE Rose
ARIMA(2,1,3)	38.817929
ARIMA(2,1,2)	38.871197

2. SARIMA Automated –

- We create a grid of all possible combinations of (p, d, q) along with Seasonal (P, D, Q) & Seasonality of 12 (for both datasets)
- Range of p = Range of q = 0 to 3, Constant d = 1
- Range of Seasonal P = Range of Seasonal Q = 0 to 3, Constant D = 1, Seasonality m = 12
- Few Examples of the grid (p, d, q) (P, D, Q, m) –

Model: (0, 1, 2)(0, 0, 2, 12)

Model: (0, 1, 3)(0, 0, 3, 12)

Model: (1, 1, 0)(1, 0, 0, 12)

Model: (1, 1, 1)(1, 0, 1, 12)

Model: (1, 1, 2)(1, 0, 2, 12)

Model: (1, 1, 3)(1, 0, 3, 12)

Model: (2, 1, 0)(2, 0, 0, 12)

Model: (2, 1, 1)(2, 0, 1, 12)

Model: (2, 1, 2)(2, 0, 2, 12)

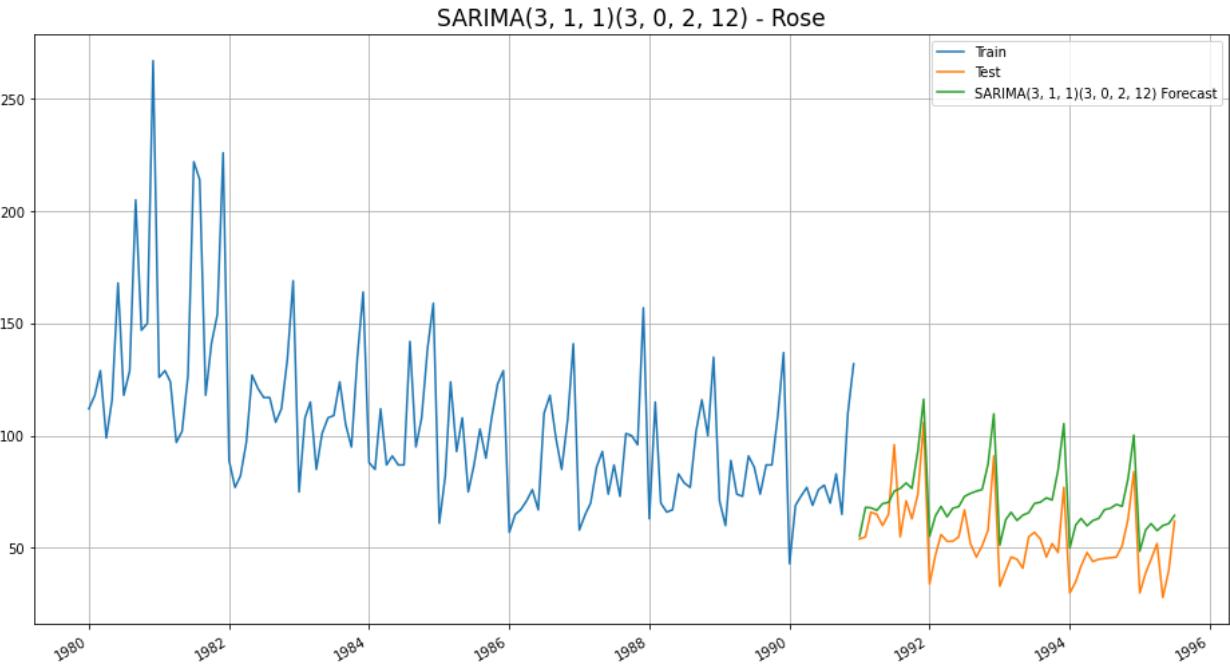
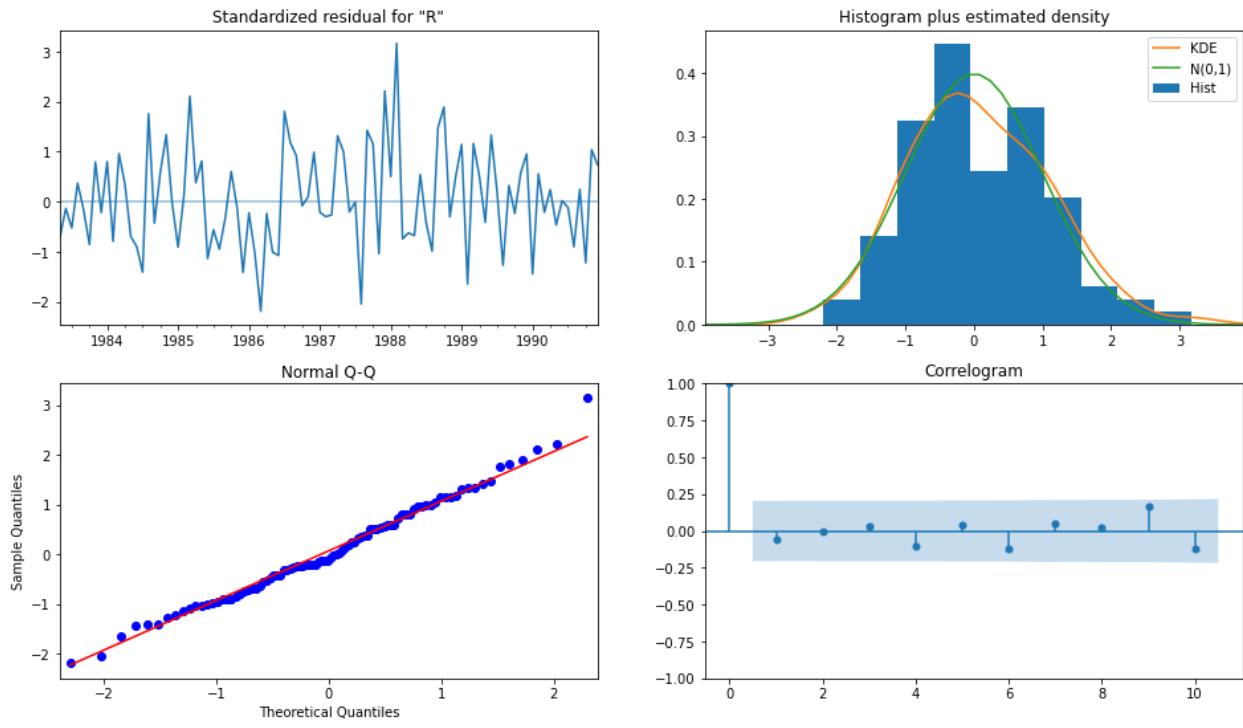
Model: (2, 1, 3)(2, 0, 3, 12)

Model: (3, 1, 0)(3, 0, 0, 12)

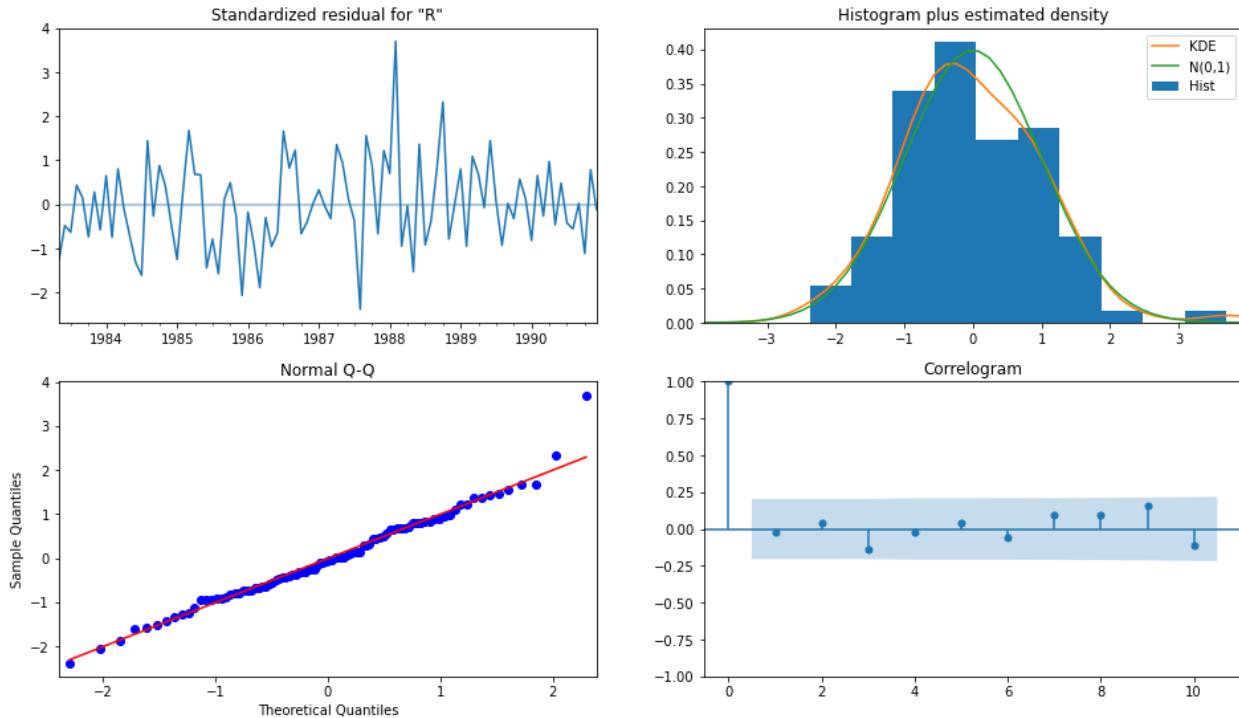
- We fit SARIMA models to each of these combinations and select with least AIC

- We fit SARIMA to this best combination of (p, d, q) (P, D, Q, m) to the Train set and forecast on the Test set. Then, we check accuracy using RMSE on Test set

- For Rose, Best Combination with Least AIC is - (3, 1, 1) (3, 0, 2, 12).



	Test RMSE Rose	Test MAPE Rose
ARIMA(2,1,3)	36.817929	75.849346
ARIMA(2,1,2)	36.871197	76.056213
SARIMA(3, 1, 1)(3, 0, 2, 12)	18.881980	36.375525



Till Now, Best Model for Rose with Least RMSE —> SARIMA (3, 1, 1) (3, 0, 2, 12)

- Till Now, Best Model for Sparkling with Least RMSE ->SARIMA (3, 1, 1) (3, 0, 0, 12)

Q 7] Build ARIMA/SARIMA models based on the cut-off points of ACF and PACF on the training data and evaluate this model on the test data using RMSE.

Auto-Correlation Function (ACF) –

- Autocorrelation refers to how correlated a time series is with its past values. e.g. with also with and so on.
- ‘Auto’ part of Autocorrelation refers to Correlation of any time instance with its previous time instance in the SAME Time Series
- ACF is the plot used to see the correlation between the points, up to and including the lag unit.
 - ACF indicates the value of ‘q’ - which is the Moving Average parameter in ARIMA / SARIMA models.

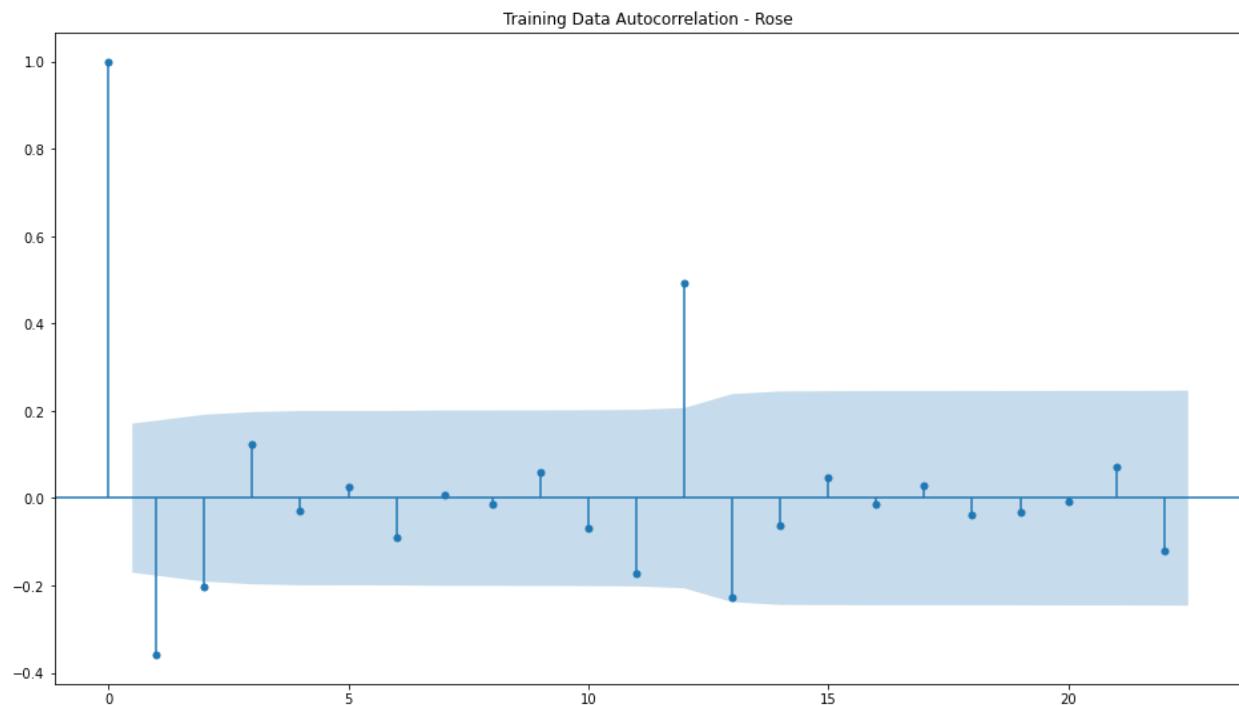
Partial Auto-Correlation Function (PACF) –

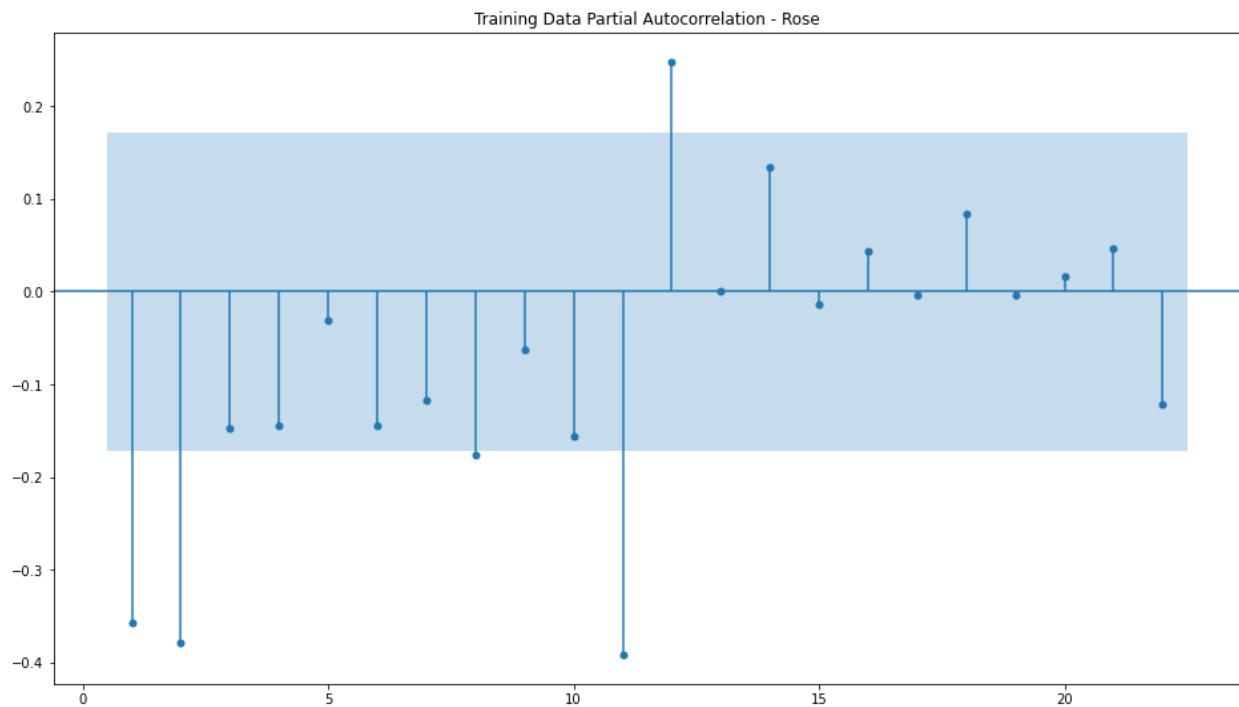
- Partial Autocorrelation refers to how correlated a time series is with its past lag values.
- For example, let lag=k, then Partial Autocorrelation is Correlation of with , ignoring the effects of all the instances between and
- PACF is the plot used to see the correlation between the lag points
- PACF indicates the value of 'p' - which is the Auto-Regressive parameter in ARIMA / SARIM A models.

ACF & PACF of Rose –

- Observing the cutoffs in ACF and PACF plots for Rose dataset, we get - FOR ARIMA —> $p = 2$, $q = 2$ and difference $d = 1$

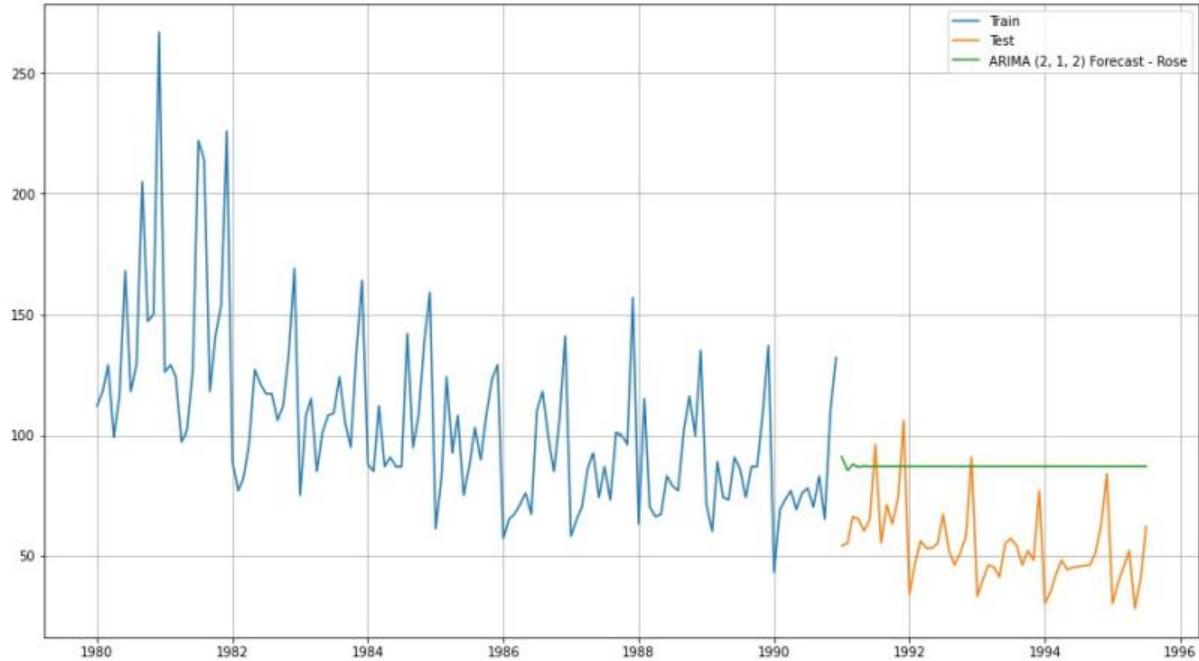
FOR SARIMA —> $p = 2$, $q = 2$, $d = 1$ and $P = 2$, $D = 1$, $Q = 2$, Seasonality=12



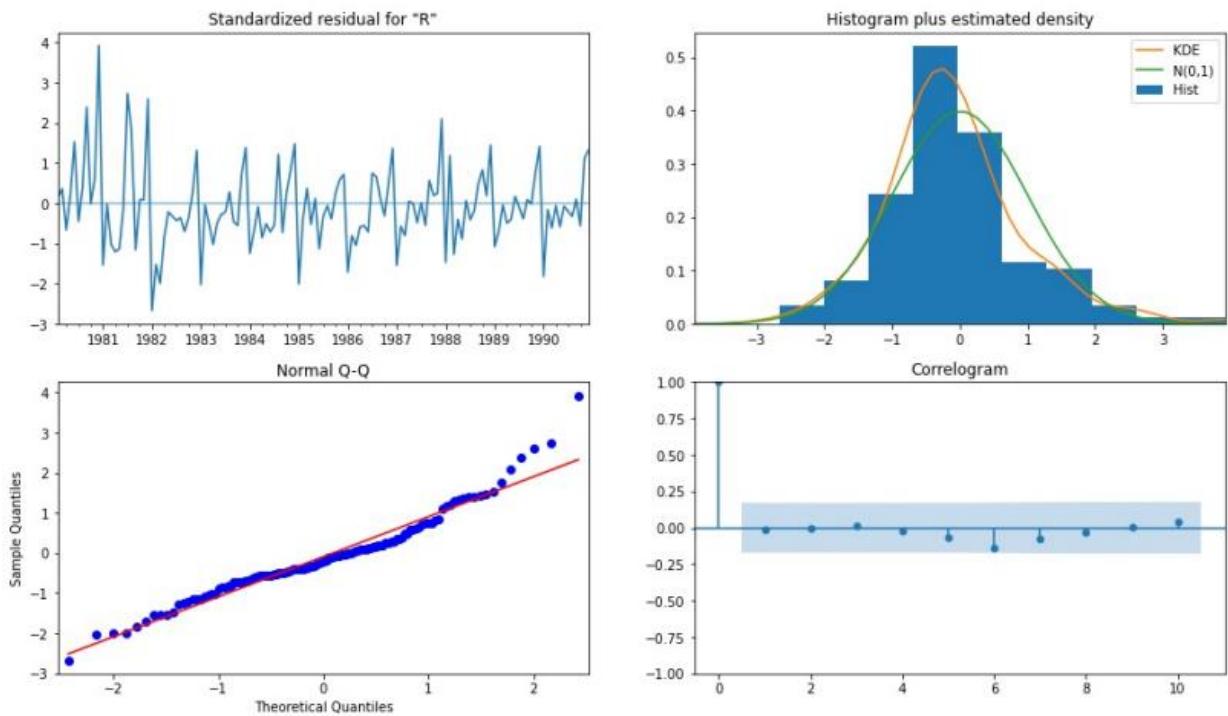


1. ARIMA Manual - Rose - (2, 1, 2)

ARIMA (2, 1, 2) Forecast - Rose



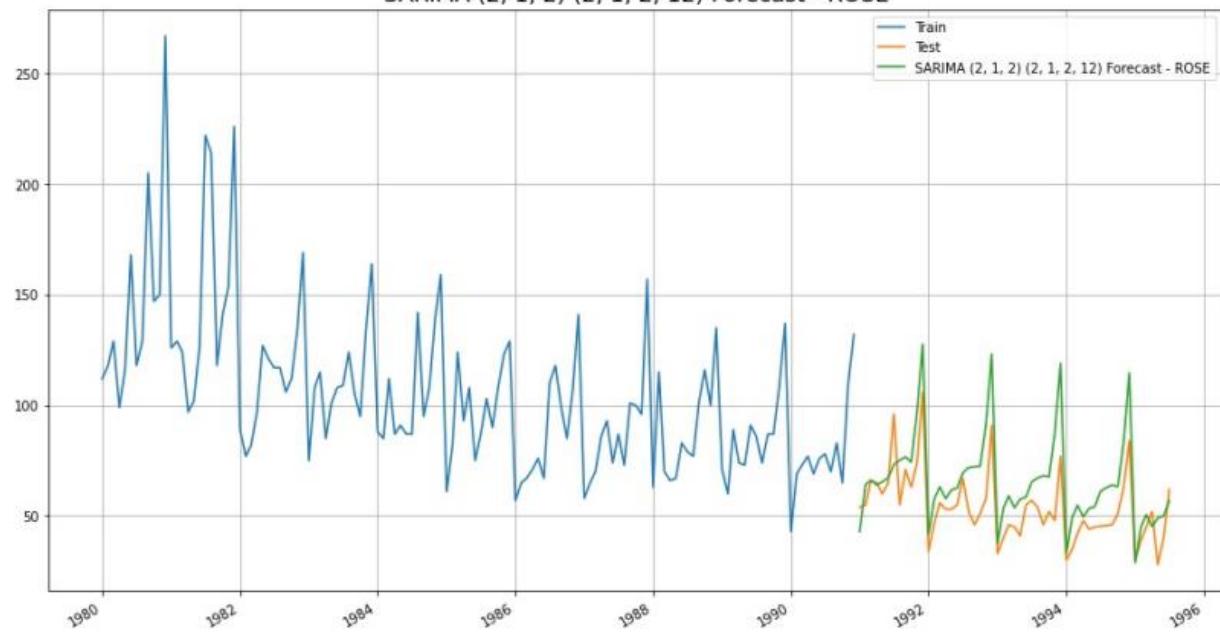
ARIMA (2, 1, 2) Diagnostic Plot - ROSE



	Test RMSE Rose	Test MAPE Rose
ARIMA(2, 1, 2)	36.87	76.06

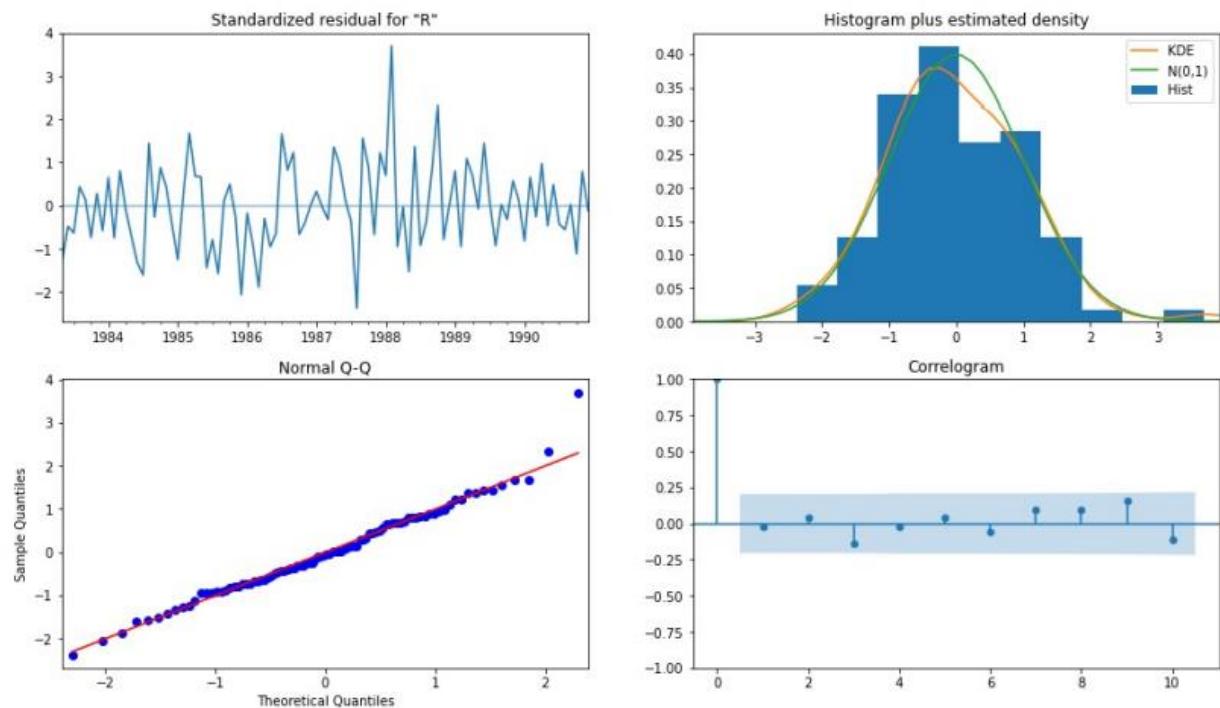
- . **SARIMA Manual - Rose - (2, 1, 2) (2, 1, 2, 12)**

SARIMA (2, 1, 2) (2, 1, 2, 12) Forecast - ROSE



https://www.kaggle.com/robinsona/rose

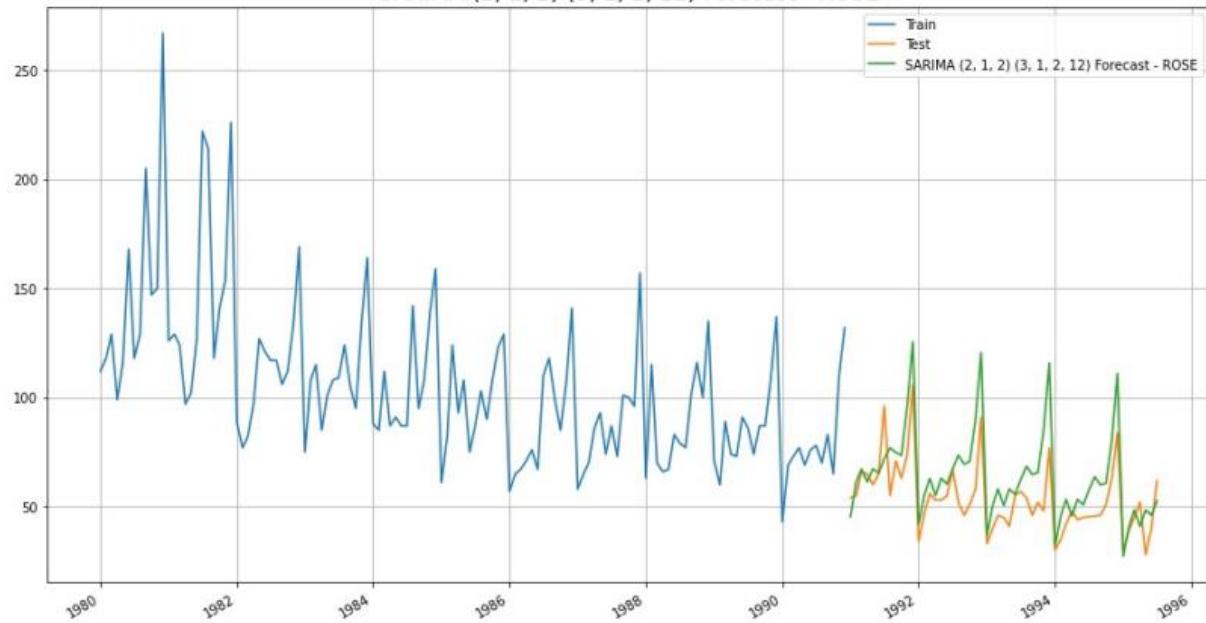
SARIMA (2, 1, 2) (2, 1, 2, 12) Diagnostic Plot - ROSE



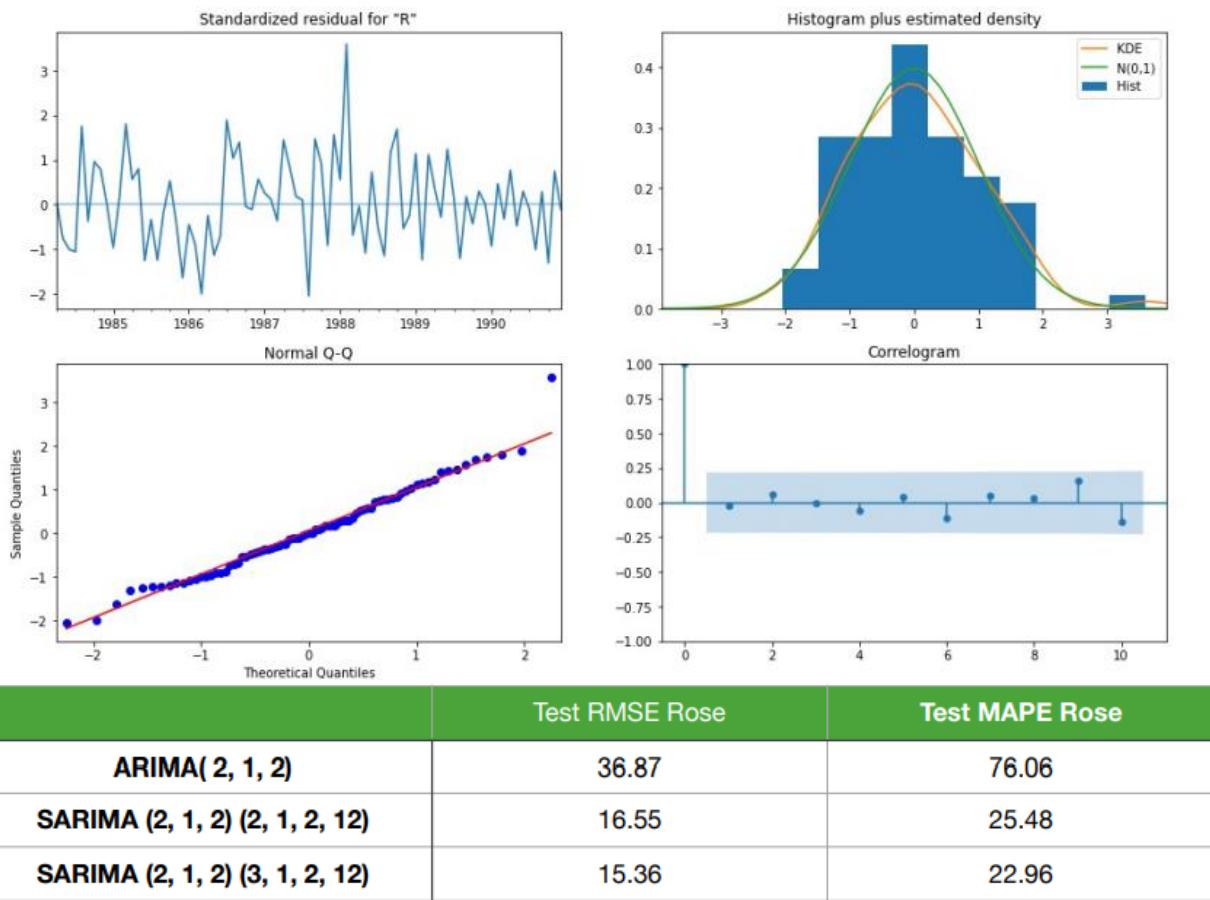
	Test RMSE Rose	Test MAPE Rose
ARIMA(2, 1, 2)	36.87	76.06
SARIMA (2, 1, 2) (2, 1, 2, 12)	16.55	25.48

SARIMA Manual - Rose - (2, 1, 2) (3, 1, 2, 12)

SARIMA (2, 1, 2) (3, 1, 2, 12) Forecast - ROSE

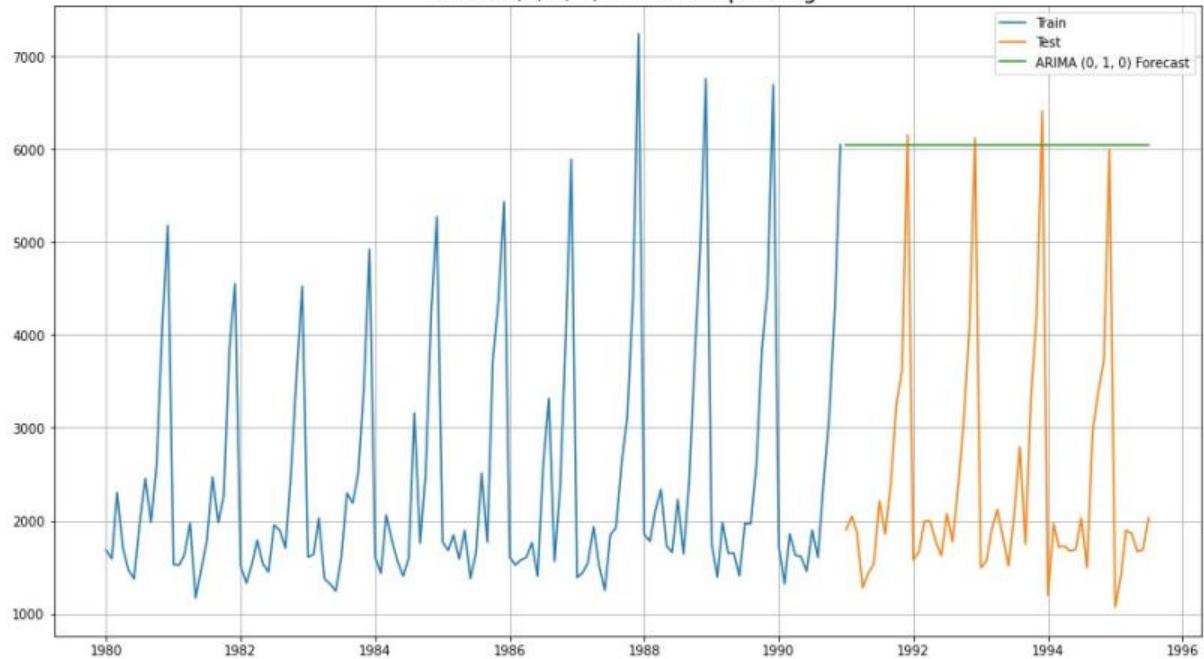


SARIMA (2, 1, 2) (3, 1, 2, 12) Diagnostic Plot - ROSE

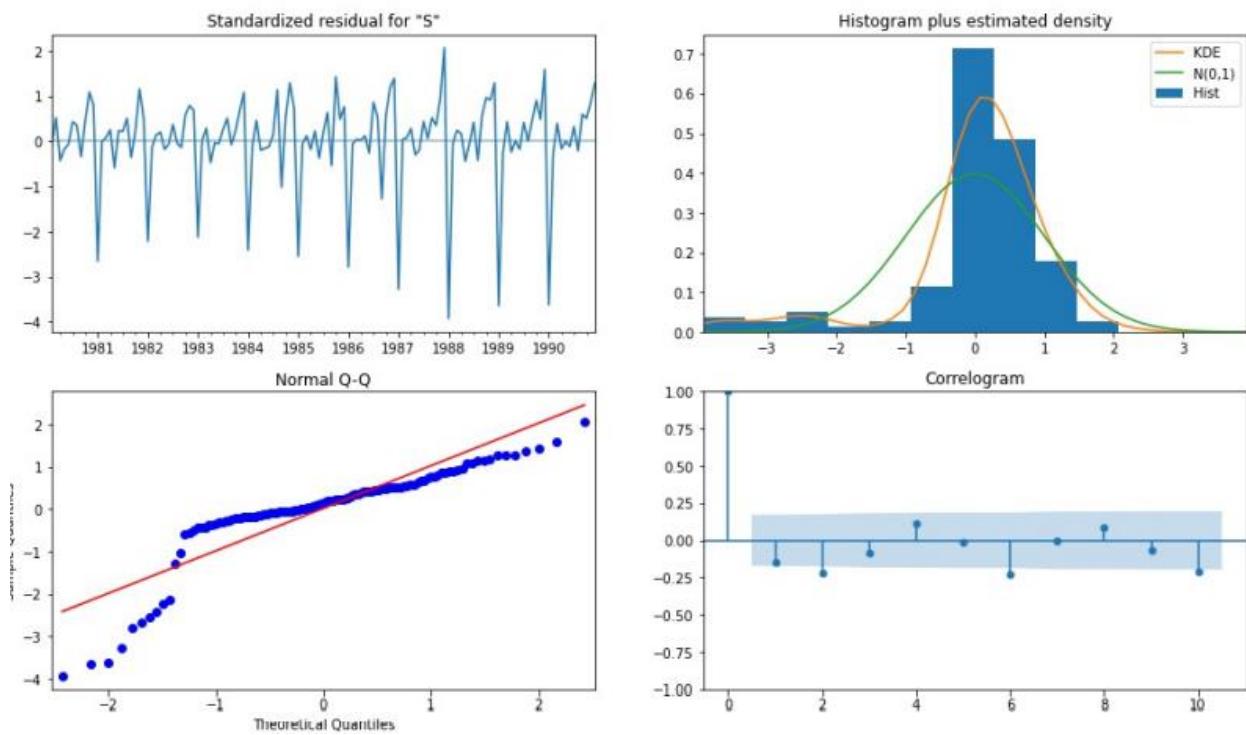


4. ARIMA Manual - Sparkling - (0, 1, 0)

ARIMA (0, 1, 0) Forecast- Sparkling



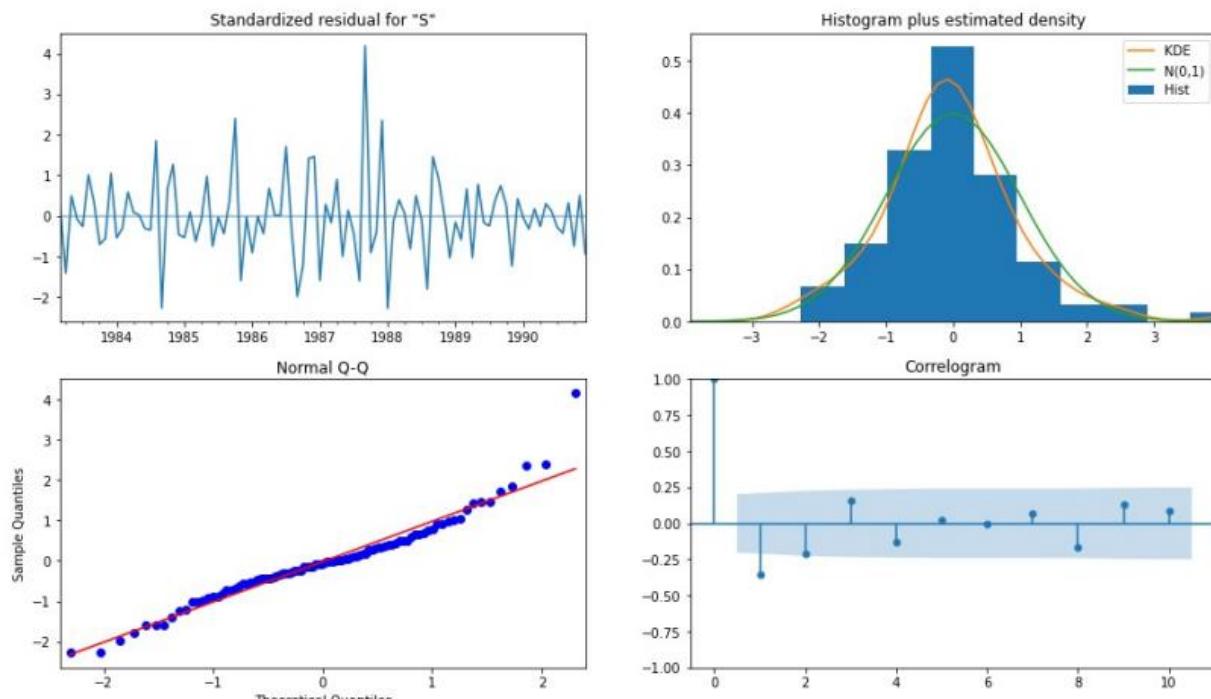
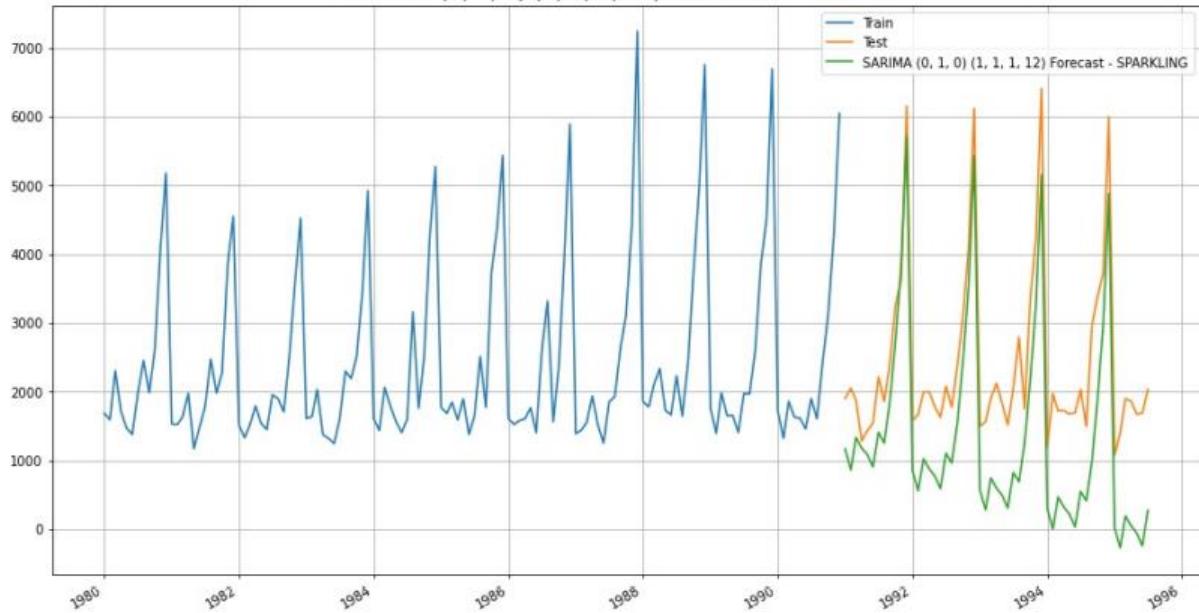
ARIMA (0, 1, 0) Diagnostic Plot - SPARKLING



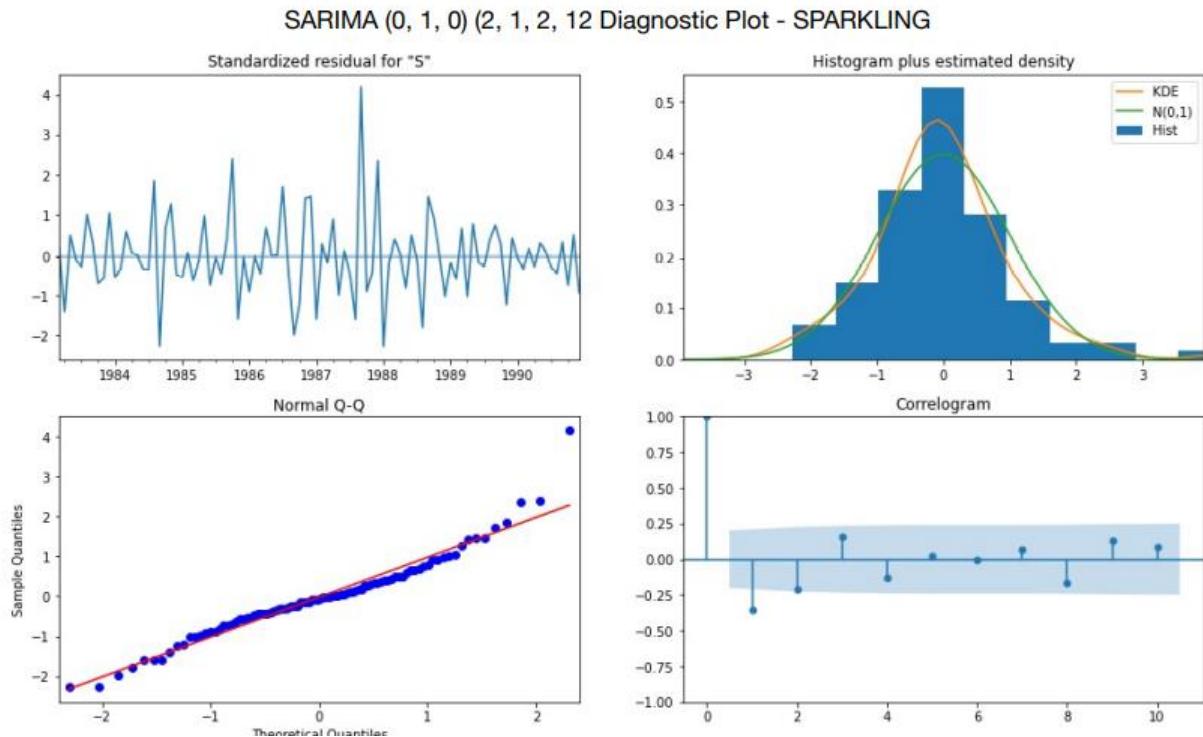
	Test RMSE Sparkling	Test MAPE Sparkling
ARIMA(0, 1, 0)	3864.28	201.33

5. SARIMA Manual - Sparkling - (0, 1, 0) (1, 1, 1, 12)

SARIMA (0, 1, 0) (1, 1, 1, 12) Forecast - SPARKLING

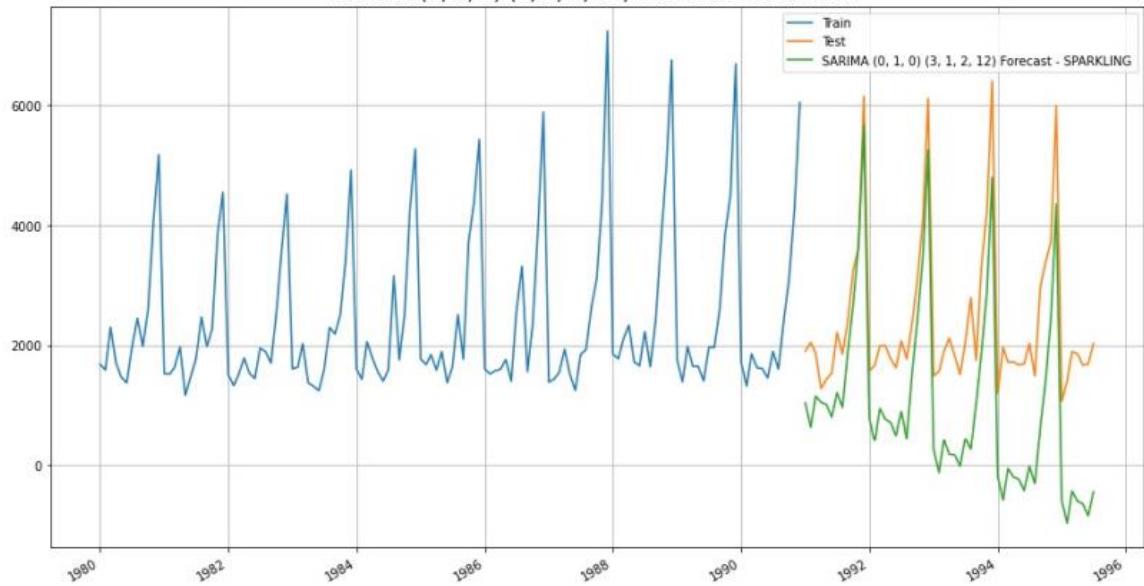


	Test RMSE Sparkling	Test MAPE Sparkling
ARIMA(0, 1, 0)	3864.28	201.33
SARIMA (0, 1, 0) (1, 1, 1, 12)	1189.84	54.87
SARIMA (0, 1, 0) (2, 1, 2, 12)	1757.73	81.79

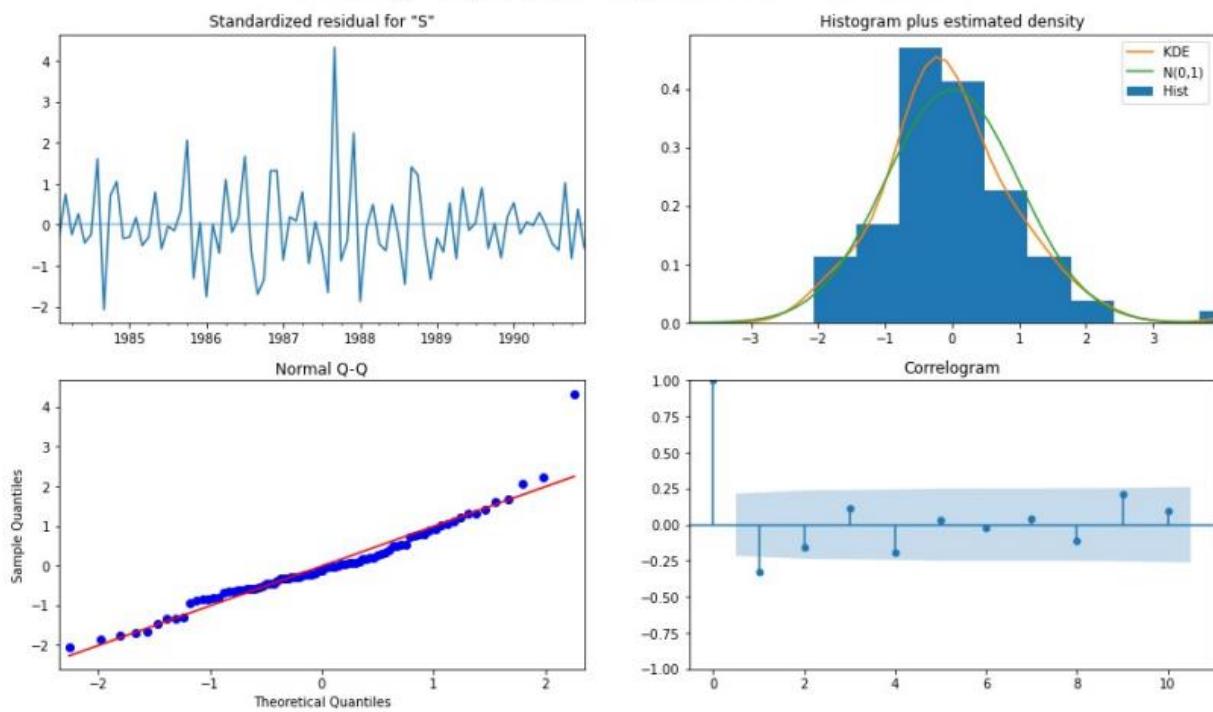


6. SARIMA Manual - Sparkling - (0, 1, 0) (2, 1, 2, 12)

SARIMA (0, 1, 0) (3, 1, 2, 12) Forecast - SPARKLING



SARIMA (0, 1, 0) (3, 1, 2, 12) Diagnostic Plot - SPARKLING



	Test RMSE Sparkling	Test MAPE Sparkling
ARIMA(0, 1, 0)	3864.28	201.33
SARIMA (0, 1, 0) (1, 1, 1, 12)	1189.84	54.87
SARIMA (0, 1, 0) (2, 1, 2, 12)	1757.73	81.79
SARIMA (0, 1, 0) (3, 1, 2, 12)	1551.65	71.57

In all Manual methods,

Best Model for Rose with Least RMSE —> SARIMA (2, 1, 2) (3, 1, 2, 12)

- In all Manual methods, Best Model for Sparkling with Least RMSE —> SARIMA (0, 1, 0) (1, 1, 1, 12)
- Seasonal P and Q - it was difficult to gauge the correct values here as the data was not enough and cutoffs were not visible
- Hence, we tried multiple combinations of Seasonal P and Q as given above

Q 8]Build a table (create a data frame) with all the models built along with their corresponding parameters and the respective RMSE values on the test data

All Models built with ROSE (sorted by RMSE) –

Model	Parameters	Test RMSE Rose
2pointTrailingMovingAverage		11.53
Triple Exponential Smoothing (Additive Season)	$\alpha = 0.0849$ $\beta = 5.52e^{-6} \approx 0.00$ $\gamma = 0.00054$	14.24
4pointTrailingMovingAverage		14.45
6pointTrailingMovingAverage		14.57
9pointTrailingMovingAverage		14.73
RegressionOnTime		15.27
Double Exponential Smoothing	$\alpha = 1.49e^{-8} \approx 0.00$ $\beta = 5.488e^{-9} \approx 0.00$	15.27
SARIMA(2,1,2)(3,1,2,12)		15.36
SARIMA(3, 1, 1)(3, 0, 2, 12)		18.88
Triple Exponential Smoothing (Multiplicative Season)	$\alpha = 0.07736$ $\beta = 0.03936$ $\gamma = 0.00083$	19.11
Triple Exponential Smoothing (Multiplicative Season, Damped Trend)	$\alpha = 0.05921$ $\beta = 0.0205$ $\gamma = 0.00405$	25.99
Triple Exponential Smoothing (Additive Season, Damped Trend)	$\alpha = 0.07842$ $\beta = 0.01153$ $\gamma = 0.07738$	26.04
Simple Exponential Smoothing	$\alpha = 0.09874$	36.80
ARIMA(2,1,3)		36.81
ARIMA(2,1,2)		36.87
SimpleAverageModel		53.46
NaiveModel		79.72

7. All Models built with SPARKLING (sorted by RMSE) –

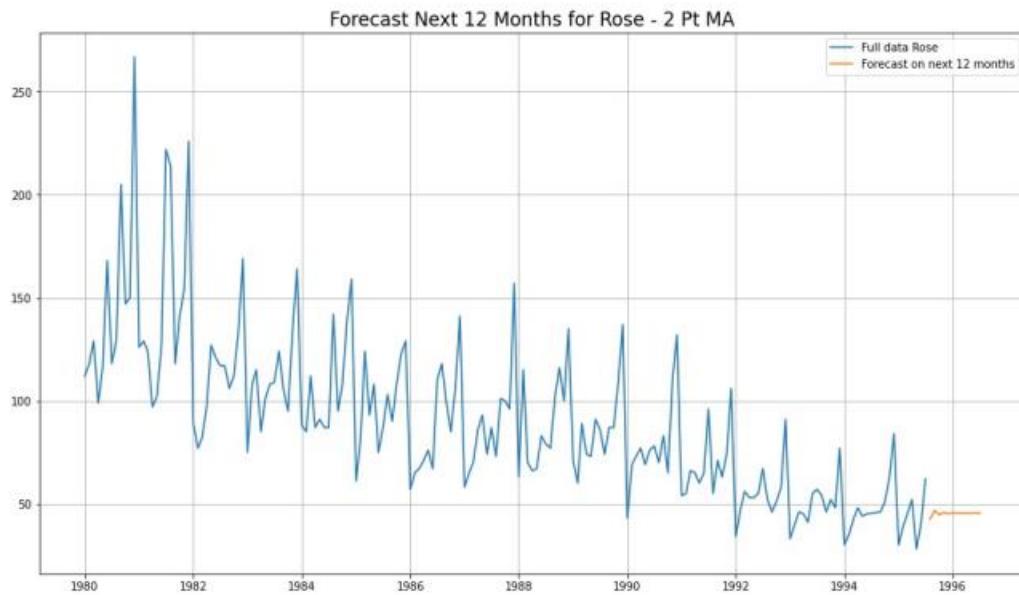
Model	Parameters	Test RMSE Sparkling
Triple Exponential Smoothing (Multiplicative Season, Damped Trend)	$\alpha = 0.11107$ $\beta = 0.03702$ $\gamma = 0.39507$	352.45
Triple Exponential Smoothing (Additive Season)	$\alpha = 0.1112$ $\beta = 0.01236$ $\gamma = 0.46071$	378.63
Triple Exponential Smoothing (Additive Season, Damped Trend)	$\alpha = 0.10062$ $\beta = 0.00018$ $\gamma = 0.51151$	378.63
Triple Exponential Smoothing (Multiplicative Season)	$\alpha = 0.11119$ $\beta = 0.04943$ $\gamma = 0.36205$	403.71
SARIMA(3,1,1)(3,0,2,12)		601.24
2pointTrailingMovingAverage		813.40
4pointTrailingMovingAverage		1156.59
SARIMA(0,1,0)(3,1,2,12)		1189.84
SimpleAverageModel		1275.08
6pointTrailingMovingAverage		1283.93
ARIMA(2,1,2)		1299.98
Simple Exponential Smoothing	$\alpha = 0.07028$	1338.00
9pointTrailingMovingAverage		1346.28
RegressionOnTime		1389.14
SARIMA(0,1,0)(3,1,2,12)		1551.65
SARIMA(0,1,0)(2,1,2,12)		1757.73
NaiveModel		3864.28
ARIMA(0,1,0)		3864.28
Double Exponential Smoothing	$\alpha = 0.665$ $\beta = 0.0001$	5291.88

Q 9]Based on the model-building exercise, build the most optimum model(s) on the complete data and predict 12 months into the future with appropriate confidence intervals/bands

- Best Models as per the Least RMSE on ROSE Test set ——>
2 Pt Trailing Moving Average - Triple Exponential Smoothing (Additive Seasonality)

- Best Model as per the Least RMSE on SPARKLING Test set —> Triple Exponential Smoothing (Multiplicative Season, Damped Trend) , , Rose Forecast Next 12 months - 2 Pt Moving Average
- This model doesn't seem to be predicting very well
- Hence, forecasting on the second best model - Triple Exponential Smoothing ETS(A, A, A) - Additive Seasonality $\alpha = 0.11107 \beta = 0.03702 \gamma = 0.39507$.

8. Rose Forecast Next 12 months - 2 Pt Moving Average

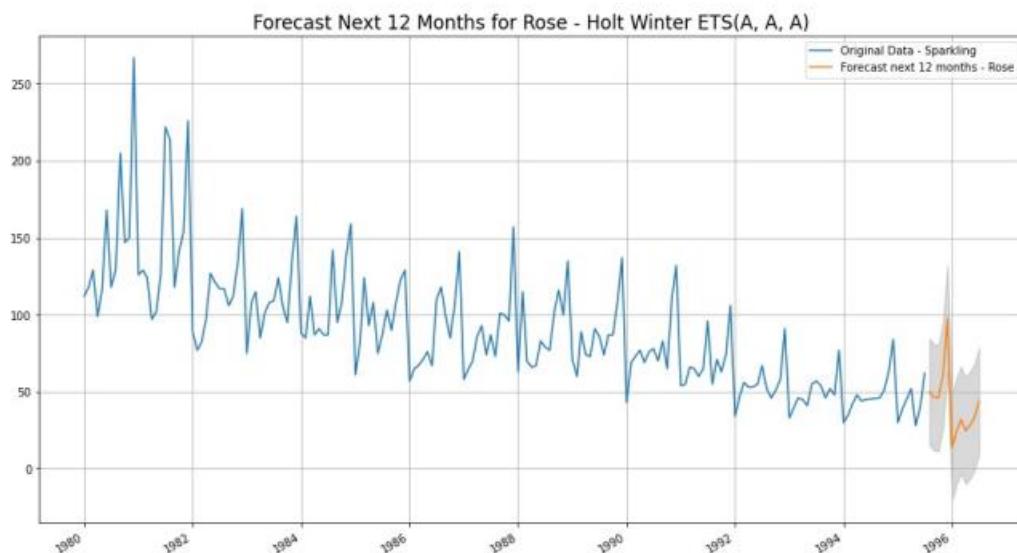


- This model doesn't seem to be predicting very well
- Hence, forecasting on the second best model - Triple Exponential Smoothing ETS(A, A, A) - Additive Seasonality.

	Forecast
1995-08-01	42.50
1995-09-01	46.75
1995-10-01	44.63
1995-11-01	45.69
1995-12-01	45.16
1996-01-01	45.42
1996-02-01	45.29
1996-03-01	45.36
1996-04-01	45.32
1996-05-01	45.34
1996-06-01	45.33
1996-07-01	45.33

Rose Forecast Values - Next 12 Months - 2 Pt Moving Average

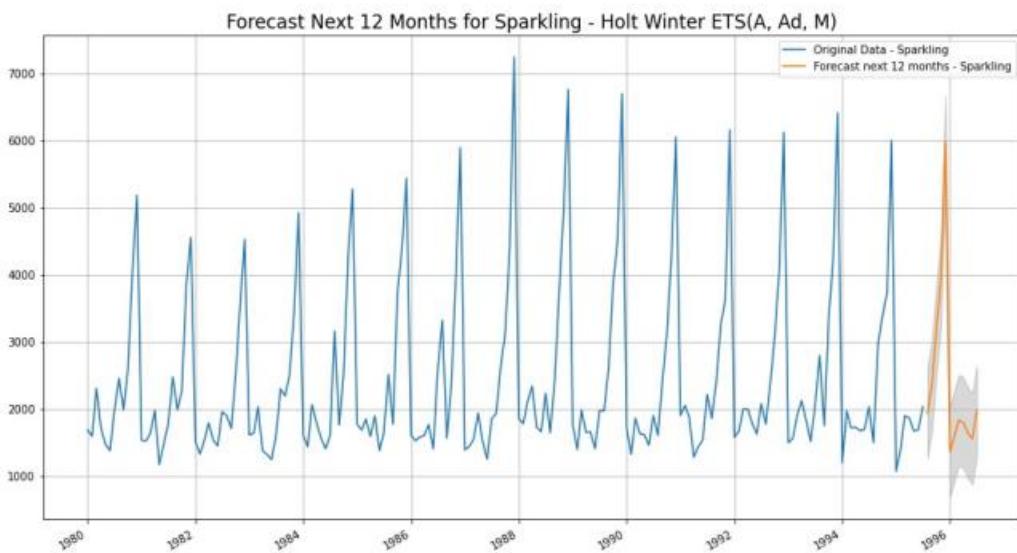
Rose Forecast Next 12 months - Triple Exponential Smoothing ETS (A, A, A)



9. Rose Forecast Values - Next 12 Months - TES - ETS(A, A, A)

	Forecast
1995-08-01	50.02
1995-09-01	46.50
1995-10-01	45.94
1995-11-01	60.72
1995-12-01	97.66
1996-01-01	13.89
1996-02-01	24.29
1996-03-01	31.91
1996-04-01	24.73
1996-05-01	28.12
1996-06-01	33.56
1996-07-01	44.09

10. Sparkling Forecast Next 12 months - Triple Exponential Smoothing ETS (A, Ad, M) - Damped Trend, Multiplicative Seasonality.



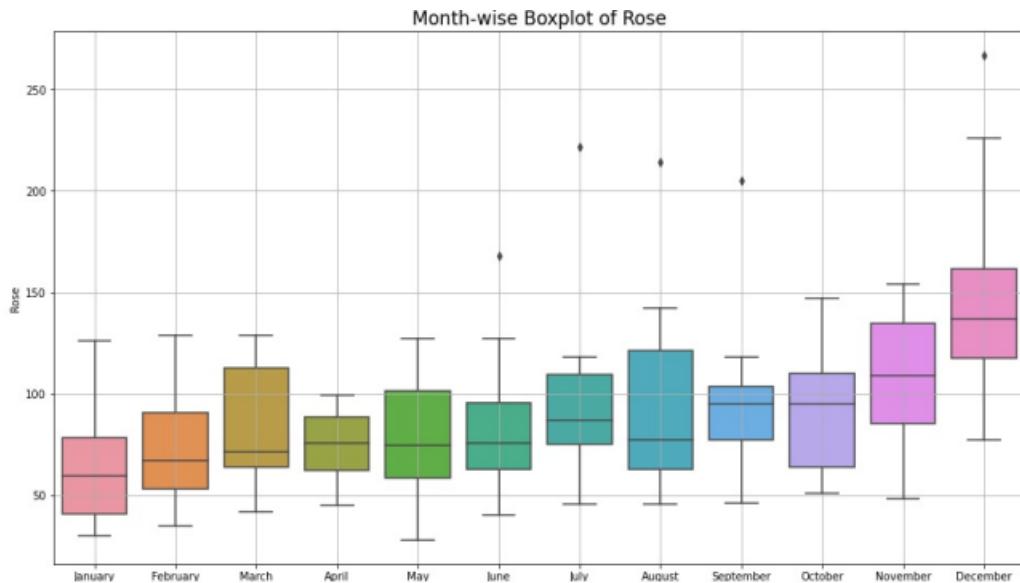
11. Sparkling Forecast Values - Next 12 Months - TES - ETS(A, Ad, M)

	Forecast
1995-08-01	1931.44
1995-09-01	2351.98
1995-10-01	3179.46
1995-11-01	3918.07
1995-12-01	5985.90
1996-01-01	1357.57
1996-02-01	1599.15
1996-03-01	1830.31
1996-04-01	1791.02
1996-05-01	1641.84
1996-06-01	1556.37
1996-07-01	1966.00

Q 10] Comment on the model thus built and report your findings and suggest the measures that the company should be taking for future sales.

12. Rose Wine Sales - Comments :

- Rose wine shows a clear trend of declining sales since 1980 - This shows decline in popularity of this variant of wine
- Also, there is a clear spike in sales seen in the last quarter of every year from Oct to Dec - This might be due to the Holiday season in this period - Highest peak in sales is seen in Dec every year
- There is also an instant crashing slump in sales in the first quarter of every year from Jan - This might be due to the after effect or hangover of Holidays.
- Sales slowly pick up only after May-June.



Rose Wine Sales - Forecast Models :

- Top 2 best models as per lowest Test RMSE were found to be - 2 Pt Moving Average and Holt-Winters - Additive Seasonality & Trend
- 2 Pt Moving Average model, when used for forecasting do not seem to give good predictions. Forecast values level out after a few iterations
- Holt-Winters seems to give a consistent forecast with respect to the data
- Hence, for final forecast of Rose Wine Sales - we choose Holt-Winters.

Rose Wine Sales - Suggestions :

- Firstly, Holiday season is around the corner and forecast shows increasing sales and sharp peak in Dec. Hence, Company should stock up
- But Declining sales of Rose Wine over the long period should be investigated with more data crunching
- Company can rebrand its Rose variant along-with a new Wine-master

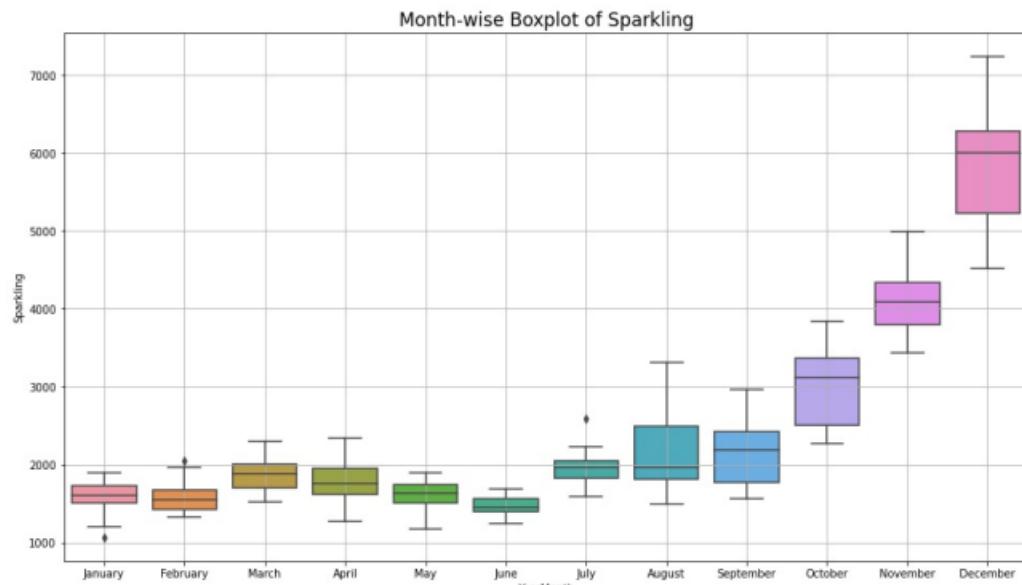
Company should take advantage of the oncoming spike from Aug-Oct by introducing aggressive offers and Ad campaigns.

- This will entice first time Wine drinkers and fence sitters (who don't have specific loyalties to any particular brand)

- Still if there is no significant upward trend in sales by this Dec, then Company has 2 options - invest in R&D or think of discontinuing this variant and come up with something completely new.

Sparkling Wine Sales - Comments :

- Sparkling wine sales don't show any upward or downward trend - This shows flat sales over long term range
- Also, there is very high spike in sales seen in the last quarter of every year from Oct to Dec
 - This might be due to the Holiday season in this period
 - Highest peak in sales is seen in Dec every year
 - Dec sales are almost 3 times of Sep sale.



Similar to Rose Wine, even in Sparkling sales, an instant crashing slump is seen in the first quarter of every year from Jan

- This might be due to the after effect or hangover of Holidays
- Sales slowly pick up only from Jul-Aug.

Sparkling Wine Sales - Forecast Models :

Triple Exponential Smoothing - Holt-Winters Models perform the best on Sparkling datasets, considering the least RMSE on Test data

- There has been incremental improvements in Test RMSE with each tuning of parameters
- Finally, for forecast of Sparkling Wine Sales
- we choose Holt-Winters with Multiplicative Seasonality and Additive Damped Trend.

Sparkling Wine Sales - Suggestions :

Even for Sparkling, Holiday season is around the corner and forecast shows increasing sales and sharp peak in Dec. Hence, Company should stock up

- Sparkling wine has great holiday sales, so this shows popularity.
- So no need to introduce any offers here but hammering Ads are suggested in these times of Oct-Dec. This will drive sales even further.
- Sparkling wines are generally associated with celebrations and mainly to burst open. •

A special designer bottle can be introduced at a cheaper price just for bursting. This will maximize profits

- Year on Year sales do not show any significant increase or decrease
- Though, Holiday spikes are extreme, but general Year on Year sales need to be investigated more. Early period from Jan should be used to do this deep Dive.